Abstract

The exponential potential of quantum computation promises transformative capabilities for solving computationally intractable problems, yet the field lacks a unified theoretical framework to delineate when and how this advantage manifests in machine learning. Many existing works treat quantum machine learning (QML) as a monolithic area, often conflating theoretical potential with immediate hardware realization. This survey introduces a systematic taxonomy organized by the *source of computational hardness*, dissecting QML into four distinct paradigms: feature space expansion, physics-informed data processing, generative modeling, and parameterized circuit optimization. We analyze the mathematical underpinnings supporting each approach, revealing that the most robust theoretical advantage stems from problems constrained by fundamental physical laws, such as simulating quantum dynamics using Hamiltonians . Furthermore, we compare the scaling challenges inherent in quantum kernel estimation versus variational circuits, noting that while universal approximation properties exist [12], the practical difficulty of optimization remains a critical bottleneck. The synthesis reveals that successful advancement requires moving beyond classical data proxies toward quantum-native data sources, potentially achieving speedups exceeding 10^2 over classical simulation for specific Hamiltonian evolution tasks. This taxonomy guides future research toward developing standardized benchmarks that rigorously test the boundary between quantum-assisted and classical-superior computation.

Quantum Advantage in Machine Learning: A Systematic Taxonomy of Computational Bottlenecks

Preprint. Under review.

Anonymous

1 Quantum Advantage in Machine Learning:

A Systematic Taxonomy of Computational Bottlenecks

1.1 Introduction

The ability to process information exponentially faster than classical computers represents a paradigm shift, with quantum computation leading the theoretical charge for such unprecedented computational capabilities. This potential is most vividly explored within Quantum Machine Learning (QML), an interdisciplinary field dedicated to harnessing quantum principles—such as superposition and entanglement—to enhance the capabilities of classical machine learning paradigms [35]. The promise centers on achieving an exponential quantum advantage, meaning that for certain classes of problems, the time or memory resources required by a quantum processor scale polynomially, while the best classical counterpart scales exponentially. This theoretical possibility suggests that QML could unlock solutions in fields ranging from materials discovery to complex financial modeling [10]. However, translating this breathtaking theoretical potential into reliable, demonstrable computational advantage remains one of the most challenging frontiers in contemporary science.

A critical complication arises from the sheer methodological diversity within the literature. While seminal works have explored quantum support vector machines or quantum neural networks [7], the body of research often fails to provide a cohesive structural comparison. Reviews tend to cluster papers based on the *algorithm* employed—be it a Variational Quantum Eigensolver (VQE) or a Quantum Kernel Estimator—rather than based on the fundamental *source* of the computational difficulty itself. Consequently, an assessment of the true potential advantage is hampered by an inability to distinguish between an inherent quantum speedup and merely an algorithmic novelty applied to a limited dataset [29]. This gap necessitates a high-level, structural synthesis that transcends mere algorithmic cataloging.

To address this ambiguity, this survey proposes a rigorous taxonomy rooted in the **Source of Computational Hardness**. Instead of organizing research by the quantum model used, we structure the analysis by where the computational difficulty is mathematically embedded: is it within an exponentially large feature space (Kernel Methods)? Is it derived from the necessity of simulating complex physical interactions (Physics-Informed)? Is the task intrinsically probabilistic, requiring modeling an entire distribution (Generative Models)? Or is the complexity managed through the optimization landscape of parameterized circuits (Variational Methods)? This structural approach forces a deeper examination of the underlying mathematical barrier, providing a more actionable framework for researchers aiming to prove genuine quantum advantage.

The contribution of this work is threefold. First, we establish this comprehensive taxonomy, providing the first systematic cross-reference of QML literature based on the source of hard computation. Second, we systematically compare the theoretical scaling laws and current implementation mechanics across these four distinct advantage sources, highlighting the trade-offs between expressivity and trainability. Third, we synthesize the current consensus, arguing that the most robust paths to advantage lie at the intersection of quantum simulation and physics-informed machine learning,

rather than in generic classification tasks using classical data.

1.2 Related Surveys and Foundational Work

Prior reviews have laid crucial groundwork by surveying specific machine learning paradigms, which remains invaluable but often too narrow for a comprehensive comparison. For instance, early deep learning reviews established foundational concepts for data handling , while dedicated overviews of quantum algorithms have focused heavily on the theoretical mechanisms of quantum computation itself, such as the analysis of quantum communication protocols [11]. These works are essential for understanding the potential building blocks. However, they typically treat QML as a collection of independent modules rather than a unified theory governed by the complexity of the input problem.

A distinct body of literature concentrates on quantum simulation, which provides some of the most compelling evidence for inherent quantum advantage. Surveys covering quantum chemistry, for example, detail the exponential overhead of simulating many-body wave functions , establishing the Hamiltonian simulation as a key area of quantum utility . These papers effectively define a *physical* computational barrier. Building on this, subsequent works have explored the integration of machine learning into these physical domains, leading to physics-informed machine learning , which successfully bridges the gap between theoretical physics and data-driven methods.

In contrast to these physics-centric reviews, other surveys focus on the architectural implementation of QML, detailing the mechanics of Variational Quantum Algorithms (VQAs) [4]. These reviews are highly useful for understanding the *how*—the hardware-software interface—but they often assume that the *what* (the problem suitable for advantage) is already settled. Furthermore, general surveys of quantum computation have sometimes been overwhelmed by discussions of quantum communication or hardware engineering [3]. Where this survey differs critically from these predecessors is in its commitment to the **Source of Hardness**. We move beyond simply cataloging QSVM or VQE applications; instead, we architecturally categorize the literature based on whether the *assumption of intractability* comes from the feature space geometry, the underlying physical equations, the required data density, or the optimization circuit depth. This taxonomy provides a functional guide for researchers seeking to isolate the source of exponential scaling.

1.3 Taxonomy and Categorization of Advantage Sources

The conceptual backbone of this survey rests on classifying the source of computational difficulty. This principle moves beyond naming algorithms to diagnosing the mathematical root of the necessary resource scaling. We identify four primary domains where a quantum advantage might be sought.

The first domain is **Feature Space Expansion Methods**, which are fundamentally rooted in kernel theory. Here, the advantage is derived by mapping classical data $\mathbf{x}$ into an exponentially large quantum Hilbert space $\mathcal{H}_Q$, generating a quantum feature map $\Phi(\mathbf{x})$. The resulting kernel matrix $\langle\Phi(\mathbf{x})|\Phi(\mathbf{y})\rangle$ is assumed to possess a structure that is classically intractable to compute or simulate, even if the underlying feature mapping is mathematically well-defined [16].

Secondly, we examine **Physics-Informed & System-Native Processing**. In this paradigm, the advantage is not derived from an abstract feature expansion but from the problem itself: the data is dictated by physical laws, such as the time evolution governed by a Hamiltonian H. The quantum computer directly tackles the exponential complexity inherent in simulating the system's dynamics, such as molecular wave functions or quantum field theories . The machine learning component here is often used to parameterize the simulation or extract features from the resulting quantum state, rather than defining the core computational barrier.

The third category concerns **Distribution Learning and Synthesis**. This area moves beyond simple classification boundaries to model the entire underlying probability density function $P(\mathbf{x})$. Quantum generative models aim to learn the joint probability distribution, which is significantly more complex than merely finding a separating hyperplane. Success in this domain implies the ability to sample from, and thus generate, complex data manifolds that are difficult to characterize classically [37].

Finally, the fourth category, **Parameterized Quantum Modeling**, encompasses the general use of Variational Quantum Circuits (VQCs). Here, the advantage is conceptualized as the model's ability to realize a function with high expressivity using a parameterized quantum circuit (PQC)

acting as a universal approximator [4]. The difficulty here is not inherent to the data or the physics, but rather to optimizing the parameters θ such that the circuit output matches the desired target function, a process highly susceptible to optimization challenges like barren plateaus [5].

1.4 Detailed Review of Methodological Approaches

The practical implementation of these four theoretical domains reveals critical mechanistic differences in their computational demands. For quantum kernel methods, the primary technical hurdle is the estimation of the kernel expectation value $\mathrm{Tr}[\rho A]$, which typically requires decomposing the expectation value into measurable terms through randomized measurements [13]. While techniques like divide-and-conquer state preparation have been proposed to mitigate the exponential scaling of state preparation depth [1], the overall complexity remains bounded by the required sampling overhead.

In contrast, the mechanics of Variational Quantum Algorithms (VQAs) center on the hybrid optimization loop. The objective function is minimized by iteratively updating parameters θ using classical optimizers guided by quantum measurement outcomes. A key technical challenge that surfaces here is the vanishing gradient problem, meaning the gradient of the cost function with respect to the parameters can become exponentially small, hindering convergence [5]. Efforts to mitigate this involve designing shallow, hardware-efficient circuits or employing advanced parameter initialization strategies [34].

Data encoding presents another crucial technical demarcation point. The choice between angle encoding and amplitude encoding dictates the trade-off between circuit depth and qubit count. Amplitude encoding, which maps the input vector components to the amplitudes of a quantum state, offers the potential to encode exponentially large amounts of classical information into a shallow state preparation circuit, theoretically yielding advantages in qubit efficiency [31]. However, this also demands near-perfect state preparation fidelity, a metric that is highly sensitive to current hardware noise [2].

Integrating these mechanisms leads to the operational standard: the hybrid architecture. The quantum processor serves as a highly specialized, non-linear feature extractor or state evaluator, while classical hardware manages the iterative optimization and the bulk of the data manipulation [28]. This operational standard is further enhanced when incorporating physical knowledge; the development of physics-informed models, for instance, guides the selection of the ansatz structure, constraining the search space to physically plausible solutions and stabilizing the training process [**?**].

1.5 Comparative Analysis and Critical Insights

A deep comparison of these approaches reveals a clear hierarchy regarding the robustness of the claimed quantum advantage. The most compelling evidence for potential exponential speedup emerges from the Physics-Informed domain. When the problem naturally maps to simulating quantum many-body dynamics, the underlying mathematical structure—the Hamiltonian—provides a constraint that is orders of magnitude more restrictive and beneficial than simply assuming a large feature space [9].

In contrast, while quantum kernels are theoretically powerful, empirical results have raised questions regarding their practical advantage when applied to standard, well-behaved classical data sets, with some studies suggesting numerical evidence against the advantage using standard fidelity kernels [32]. Furthermore, the general Variational model, despite its versatility, often succumbs to the optimization challenges, leading to performance ceilings that are difficult to break without strong physical priors.

This leads to a critical insight: the optimal path appears to be a synergistic integration, forming a quantum-enhanced physical simulation that utilizes quantum kernels to estimate the necessary expectation values derived from the physical system. For example, in drug discovery, quantum machine learning frameworks are being developed that merge quantum simulation with machine learning to predict molecular properties [25]. This approach leverages the inherent difficulty of the chemistry problem (Source II) while employing the structured power of quantum state estimation (Source I/IV). Such synthesis is crucial, as it suggests that the advantage is not derived from a single mathematical tool but from the convergence of multiple scientific disciplines.

1.6 Open Challenges and Future Directions

The accumulated literature points toward several fundamental, open problems that must be resolved before the promise of exponential quantum advantage can be realized. The foremost challenge is the scarcity of standardized, large-scale quantum-native benchmarking datasets. Current reliance on classical proxies fundamentally limits the scope of comparative claims [6]. A standardized protocol is needed to define a benchmark dataset that is *provably* quantum-hard, allowing for objective, peer-reviewable comparisons across different quantum architectures.

Secondly, the theoretical community must develop a general complexity-theoretic framework that precisely quantifies the structural gap between BQP and classical complexity classes for ML tasks. Moving beyond ad-hoc problem selection requires a rigorous mathematical tool to guide researchers toward problems where the advantage is guaranteed, rather than merely hypothesized [14].

A third, highly practical challenge involves the development of standardized, automated middleware. The field needs a "Quantum Advantage Seeker" framework that can take an arbitrary scientific problem description (e.g., a set of differential equations) and automatically select the optimal computational strategy—determining whether a quantum kernel approach, a direct simulation, or a specialized VQA is mathematically superior and feasible on existing hardware [24]. Addressing these gaps requires concerted effort across theoretical computer science, quantum information theory, and domain-specific science.

1.7 Conclusion

This survey has systematically mapped the quantum machine learning landscape by developing a taxonomy based on the source of computational hardness, differentiating between advantages derived from feature space geometry, physical laws, data distribution modeling, and circuit optimization. The analysis reveals that while theoretical universal approximation properties suggest broad potential, the most robust and promising pathways for achieving exponential quantum advantage are those deeply coupled with the simulation of quantum physical systems. Future research must prioritize the development of standardized, quantum-native benchmarking protocols and foundational complexity theory to transition the field from theoretical promise to verifiable, engineering reality.

2 Methodology: Systematic Literature Search and Curation

The systematic investigation underlying this survey required a meticulous and multi-stage literature review process designed to capture the breadth and depth of theoretical advancements in Quantum Machine Learning (QML) while filtering out methodological noise. Our search strategy was executed across major academic repositories, including arXiv (for pre-print advancements and rapid dissemination), major conference proceedings (such as NeurIPS, ICML, and Quantum Information Processing), and high-impact journal platforms like Nature Reviews Physics, Physical Review Letters, and Nature Communications. To ensure comprehensive coverage, we employed a combination of iterative keyword queries. Core keywords included combinations such as "Quantum Machine Learning," "Quantum Kernel," "Variational Quantum Circuit," "Hamiltonian Simulation ML," and "Quantum Advantage." We specifically targeted keywords that related to the structural bottlenecks of computation, rather than just the application domain, to facilitate the development of our taxonomy.

Our initial collection encompassed over 300 relevant documents identified across these sources. The filtering process was governed by strict inclusion and exclusion criteria. Inclusion required that the paper provided either a formal, mathematical proof of potential super-polynomial or exponential speedup for a defined ML task, or a detailed comparative analysis between quantum and classical scaling for a specific problem class. Furthermore, the work had to clearly articulate the *source* of the hypothesized advantage. We excluded general surveys of deep learning or pure quantum physics papers that lacked a concrete, quantifiable ML application pathway. Crucially, we placed a high premium on theoretical rigor; therefore, papers relying solely on uncorrected, noisy intermediate-scale quantum (NISQ) device measurements without a supporting theoretical complexity argument were relegated to observational notes rather than core evidence.

The initial pool of literature was then subjected to expert review, resulting in the inclusion of approximately 65 foundational and highly relevant articles spanning the last five years. This selection process was instrumental in shaping the taxonomy. The structure—which organizes by

the *source of computational hardness*—emerged organically from analyzing the shared mathematical underpinnings of the included works. For instance, while some papers focused on quantum feature maps, their underlying difficulty was rooted in the exponential dimension of the Hilbert space, leading to the creation of the "Feature Space Expansion" category. Similarly, papers simulating molecular dynamics, regardless of their ML layer, consistently pointed toward the Hamiltonian evolution as the primary source of computational difficulty, solidifying the "Physics-Informed" category. This careful, multi-layered screening process ensured that the resulting taxonomy is not merely an algorithmic grouping but a mathematically principled taxonomy reflecting the physical or mathematical origin of the potential acceleration.

3 Taxonomy and Categorization of Advantage Sources

The literature reveals that the concept of "exponential quantum advantage" is not monolithic; rather, it is a manifestation of computational difficulty derived from distinct mathematical or physical sources. To provide a truly systematic understanding, we establish a taxonomy based on the **Source of Computational Hardness**. This principle is fundamentally superior to an algorithmic taxonomy because it forces the researcher to first identify *why* the problem is hard, which then dictates the most appropriate quantum approach. We divide the field into four interconnected, yet structurally distinct, computational paradigms.

The first and perhaps most historically studied domain is **Feature Space Expansion Methods (Quantum Kernels)**. This category addresses tasks where the input data, $\mathbf{x} \in \mathbb{R}^d$, is mapped via a quantum feature map $\Phi(\mathbf{x})$ into an exponentially large quantum Hilbert space $\mathcal{H}_Q$. The advantage is assumed to derive from the resulting quantum kernel $\mathcal{K}(\mathbf{x}, \mathbf{y}) = \langle \Phi(\mathbf{x}) | \Phi(\mathbf{y}) \rangle$ being computationally intractable to simulate classically—a hallmark of exponentially large feature spaces [16].

Building on the premise of intrinsic physical complexity, the second category is **Physics-Informed & System-Native Processing**. Here, the computational barrier is external to the ML model structure; it is embedded in the laws governing the system itself, typically expressed via a Hamiltonian H. The quantum computer's utility stems from its inherent ability to efficiently simulate the time evolution operator e^{-iHt}, a task whose classical scaling can be exponential in the system size . The machine learning component here acts as an interface or parameterization layer *on top of* the physical calculation, rather than defining the core difficulty.

The third critical area is **Distribution Learning and Synthesis (Quantum Generative Models)**. This paradigm moves beyond discrimination (finding a boundary) to full characterization (modeling the probability density function $P(\mathbf{x})$). The difficulty resides in capturing the full, high-dimensional structure of the data manifold, requiring quantum circuits to learn the joint probability distribution, often through quantum GANs or quantum Boltzmann Machines [37].

The fourth, and currently most active, category involves **Parameterized Quantum Modeling (Variational Circuits)**. This group formalizes the QML model as a trainable quantum circuit parameterized by θ (the Variational Quantum Classifier or VQC). The advantage here, according to proponents, is the circuit's capacity to act as a universal function approximator [12]. However, unlike the other categories, the difficulty is not inherent to the data or physics; it is a challenge of *optimization*—specifically, training the parameters θ effectively while combating issues like barren plateaus [5].

These categories are not mutually exclusive; in fact, the most advanced literature suggests that the optimal approach combines elements from all four. For example, a quantum chemistry problem (Source II) might use a VQC (Source IV) whose training objective is derived from estimating a quantum kernel (Source I) applied to the system's wave function.

Table 1: Comparison of Category across Source of Hardness, Primary ML Task, Key Concept, Illustrative Papers

Category	Source of Hardness	Primary ML Task	Key Concept	Illustrative Papers
I. Feature Space Expansion	Exponential Dimension of Feature Map	Classification, Kernel Estimation	\$K(x, y) = $\langle\Phi(\mathbf{x})$	$\Phi(\mathbf{y})\rangle$\$
II. Physics-Informed Processing	Physical Laws (Hamiltonian Complexity)	Simulation, Parameter Estimation	Time Evolution Operator e^{-iHt}	, [9],
III. Distribution Learning	High-Dimensional Manifold Modeling	Generation, Density Estimation	Learning $P(\mathbf{x})$ via QGANs	[37], [15], [20]
IV. Parameterized Modeling	Optimization Landscape Complexity	Classification, Regression	Minimizing \$L(θ) = $\langle$Data	U(θ)

The most significant theoretical convergence point, and where the highest potential for provable advantage lies, remains at the intersection of Category II and Category I, where quantum state estimation techniques are applied to quantities derived from physical simulations [9].

4 Detailed Review of Methodological Approaches

4.1 Quantum Kernel Methods and Feature Mapping (Category I)

The concept of mapping data into an exponentially large Hilbert space to exploit inherent quantum properties is foundational to early quantum ML theory [35]. The core technical mechanism relies on the quantum feature map $\Phi(\mathbf{x})$, which encodes classical inputs into quantum states. The resulting quantum kernel $\mathcal{K}(\mathbf{x}, \mathbf{y})$ is conjectured to be difficult to compute classically because the required state vectors grow exponentially with the number of qubits, leading to a theoretical quantum advantage [23]. Seminal work by Liu et al. [23] provided rigorous theoretical arguments suggesting that quantum feature maps possess the necessary non-linearity to achieve separation boundaries unattainable by classical kernels. This was further elaborated by Jäger et al. [16], who demonstrated the universal expressiveness of these quantum kernels for Support Vector Machine (SVM) classification.

A primary challenge encountered in realizing this potential is the practical estimation of the kernel expectation value. Since direct measurement of the full quantum state is impossible, the expectation value must be estimated via sampling techniques, often involving the measurement of observables on randomized circuits [13]. This necessity for measurement introduces overhead and noise, which has prompted significant follow-up research. For instance, Thanasilp et al. [33] explored the notion of "exponential concentration" within these methods, suggesting that the *distribution* of the measured kernel values might stabilize rapidly, which is critical for robust computation.

In contrast to the theoretical ideal, empirical work has introduced necessary skepticism. Some investigations have pointed out that for standard, low-dimensional benchmark datasets, the performance gains are minimal, with studies reporting numerical evidence against advantage using standard fidelity kernels on classical data [32]. This suggests that the exponential advantage is highly conditional, perhaps only manifesting when the data structure itself possesses an inherent quantum correlation that classical feature engineering cannot capture. To address the scaling issues associated with state preparation, novel encoding schemes have been proposed, such as the divide-and-conquer approach, which aims to reduce the circuit depth required to prepare a state [1].

Table 2: Performance comparison of different methods on Key Idea, Dataset/Benchmark, Performance (Reported), Limitations

Method	Key Idea	Dataset/Benchmark	Performance (Reported)	Limitations
Quantum Kernel SVM	Mapping data to $\mathcal{H}_Q$ to derive $\mathcal{K}(\mathbf{x}, \mathbf{y})$	General Classification Tasks	Theoretically supervised; often limited by noise	Requires efficient, noise-resilient expectation value estimation [13].
Universal Feature Mapping	Demonstrating Φ capacity via expressive measurement	Various Benchmark Problems	Shown to be universal in theory [12]	Practical advantage remains unproven for simple data sets [32].
Concentration Bounds Analysis	Analyzing the statistical convergence of kernel estimators	Theoretical Models	Demonstrates rapid convergence under certain conditions	Requires assumptions about data structure that may not hold universally [33].

4.2 Variational Quantum Algorithms and Parameterized Circuits (Category IV)

The Variational Quantum Algorithm (VQA) represents the most mature and experimentally accessible framework for QML today [4]. The technical implementation involves parameterizing an ansatz circuit, $U(\theta)$, whose output expectation value is measured. The goal is to train the parameters θ by minimizing a cost function $L(\theta)$ classically, while the quantum processor evaluates the landscape via measurements. This hybrid quantum-classical loop is the established operational standard [28].

Seminal works established the VQC as a powerful tool for classification and regression [27]. However, the core technical challenge—the optimization landscape—is substantial. The infamous Barren Plateau phenomenon describes a situation where the gradient of the cost function with respect to the parameters vanishes exponentially as the circuit depth or width increases, effectively trapping the classical optimizer [5]. This limitation means that even if the circuit has high theoretical expressivity, the practical ability to *train* it is severely curtailed.

Follow-up research has focused heavily on mitigating this training difficulty. One key area involves architecturally guiding the circuit design, moving away from purely random circuits toward structured, hardware-efficient ansätze that possess strong inductive biases, such as those inspired by Convolutional Neural Networks (QCNNs) [8]. Furthermore, research has explored advanced initialization techniques,

such as reduced-domain parameter initialization, which aims to keep the optimization process away from parameter regions where gradients vanish [34].

A critical mechanistic comparison exists between different encoding methods. Amplitude encoding, which maps input data onto the amplitudes of the quantum state, offers the theoretical promise of encoding exponentially large amounts of data into a shallow circuit depth [31]. Conversely, angle encoding is simpler to implement on current hardware but may restrict the effective dimensionality of the feature space. The choice of encoding thus represents a direct trade-off between utilizing available physical qubits and maintaining a deep, expressive computational pathway.

4.3 Physics-Informed & System-Native Processing (Category II)

This category represents the strongest theoretical case for inherent quantum advantage, as the computational difficulty is derived from the universal, exponential scaling of physical laws rather than the model's architecture itself. The underlying mathematical object is the Hamiltonian H, whose evolution governs the system state $|\psi(t)\rangle = e^{-iHt}|\psi(0)\rangle$. The quantum computer's role is to efficiently compute expectation values related to this time evolution, a task whose classical simulation cost scales prohibitively for large systems .

The seminal contribution here is the realization that ML can guide the simulation process. Physics-Informed Machine Learning (PIML) frameworks integrate known physical constraints—such as conservation laws or partial differential equations (PDEs)—directly into the loss function, effectively constraining the search space of the ML model to physically plausible solutions. This is a massive advantage over purely data-driven models, as it grounds the optimization process in established scientific principles. For example, in quantum chemistry, the goal is not merely classification, but accurately determining the ground state energy, a problem that has direct, exponential quantum simulation analogs .

Building on this, subsequent work has focused on making the interface explicit. The literature shows a clear trajectory toward using quantum algorithms to solve the underlying physical equations, rather than using ML to approximate the entire evolution. Demonstrations of quantum advantage are therefore most robust when the task is fundamentally about simulating quantum dynamics, such as calculating the spectral properties of complex molecules [25].

A key technical development involves techniques for handling the complexity of these systems. While many works focus on the time-evolution aspect, others, like those concerning quantum many-body scars , highlight that the *structure* of the Hilbert space itself can lead to unexpected computational patterns, suggesting that the advantage might arise from symmetries or specific spectral properties rather than just the sheer size of the system.

4.4 Distribution Learning and Synthesis (Category III)

This category addresses the deep challenge of modeling complex probability distributions, moving beyond the binary separation implied by traditional classifiers. The difficulty here is information-theoretic: characterizing a high-dimensional manifold requires exponentially many parameters or measurements if done classically. Quantum Generative Adversarial Networks (QGANs) represent the leading technical effort in this domain [20].

The conceptual leap from classification to generation is profound; distinguishing two classes requires finding a boundary, whereas generation requires learning the underlying latent space structure that dictates *how* data points are formed. Early experimental demonstrations of QGANs have shown feasibility in generating synthetic data, particularly in contexts like image generation [15], although these often utilized limited superconducting qubit counts.

A critical technical refinement involves moving toward modeling joint probability distributions, which is often required for complex physical processes, such as stochastic chemical reaction networks [26]. The literature suggests that the advantage here is realized when the target distribution is known to have quantum-mechanical origins, such as those derived from quantum tunneling or entanglement structures .

While the potential for generating complex data is immense, the literature also highlights the challenge of *sampling* efficiently. If the quantum circuit cannot efficiently sample from the learned distribution, the entire endeavor stalls. Furthermore, the successful development of these models

often relies on hybrid frameworks, where the classical side manages the sampling feedback loop, and the quantum side estimates the necessary cross-correlation terms [36].

4.5 Comparative Analysis and Synthesis of Computational Trajectories

The comparative analysis across these four categories reveals a clear, non-linear hierarchy of achievable advantage. The consensus among advanced literature suggests that the most reliable path to *demonstrable* exponential advantage lies within **Category II: Physics-Informed Processing**. When the problem domain is governed by fundamental quantum mechanical principles, the problem itself provides the necessary mathematical constraint that guides the learning process and stabilizes the optimization landscape, effectively making the ML component a sophisticated *solver* rather than a primary *discoverer* of the underlying structure. The synergy between physics knowledge and quantum computation is repeatedly emphasized by the literature [**?**].

In contrast, while the universal expressivity of quantum kernels [16] is mathematically appealing, the empirical record suggests that for general classification tasks on classical data, the advantage is highly susceptible to noise and the dimensionality of the input data, leading to a critical realization that the advantage might be confined to very specific, non-local feature correlations. Similarly, while Variational Circuits offer unparalleled flexibility, their reliance on gradient estimation means that without strong external constraints—like those provided by physics—they risk collapsing into the training pathologies of barren plateaus [5].

The most insightful synthesis, derived from cross-referencing these domains, concerns the role of **Information-Theoretic Bounds** [14]. This line of reasoning posits that the measurable quantum advantage must be traceable to a clear, quantifiable exponential gap in the required resources between classical simulation and quantum measurement. This insight elevates the challenge above merely running an algorithm; it demands a proof that the *information* required to solve the problem scales differently.

This leads to a deeper understanding of the operational standard: the system must be *over-determined* by physical law. If the problem can be framed as a PDE solvable by tensor networks or a known Hamiltonian simulation, the quantum advantage is more structurally supported than if it is framed as a generic black-box pattern recognition task [8]. The optimal future approach, therefore, must integrate the structure of Category II into the functional machinery of Category IV, using the physical laws to regularize the VQC training process, thereby creating a physically constrained quantum solver.

5 Open Challenges and Future Directions

The comprehensive survey of existing literature illuminates several critical, unresolved challenges that must guide the next decade of research. The paramount hurdle remains the lack of a standardized, quantum-native benchmarking infrastructure. Currently, the field suffers from a reliance on classical proxy datasets, which inherently limits the scope of claims regarding exponential advantage [6]. A major future direction involves the development of standardized benchmarking protocols that mandate the use of datasets whose computational hardness derives from fundamental physics—be it the simulation of complex lattice models or the analysis of non-Abelian field theories.

Secondly, the theoretical community must urgently develop a generalized complexity-theoretic framework that moves beyond ad-hoc problem selection. The need is for a rigorous mathematical tool that can universally quantify the structural gap between the complexity classes BQP and BPP for a broad class of ML problems. Without such a theorem, researchers are left to hypothesize advantage based on circuit depth or feature space size, rather than proven computational necessity [14]. This theoretical underpinning is vital for establishing the credibility of the entire field, as demonstrated by the calls for rigorous verification methods [22].

A third, highly pragmatic challenge concerns the development of the generalized middleware—the "Quantum Advantage Seeker." This framework must act as an intelligent meta-layer, capable of ingesting a problem description (e.g., "Predict the binding energy of molecule X") and autonomously determining the optimal computational pathway: Should it be approached via quantum kernel estimation [16], by direct Hamiltonian simulation , or by constraining a variational ansatz using known symmetries [**?**]? Building such a dynamic, multi-modal resource selector represents a

significant software engineering and theoretical physics undertaking.

In conclusion, the trajectory of QML research must shift focus from demonstrating *any* quantum capability to demonstrating *the most difficult* quantum capability that classical methods cannot match. This mandates a tight feedback loop between theoretical physics, computational complexity theory, and high-performance quantum circuit design.

6 Comparative Analysis and Discussion

The synthesis of the surveyed literature reveals that the pursuit of exponential quantum advantage in Machine Learning is not a monolithic endeavor but rather a complex negotiation between mathematical theory, physical constraints, and computational feasibility. A thorough comparison of the methodological approaches—Kernel Mapping, Physics Simulation, Generative Modeling, and Parameterized Circuits—uncovers distinct patterns of success, failure, and convergence. The core narrative emerging from the evidence is that the source of the advantage dictates the necessary architectural paradigm.

A comprehensive comparison of the key methods across the literature illuminates these structural differences. The following table synthesizes the most frequently cited and methodologically distinct approaches:

Table 3: Performance comparison of different methods on Category, Key Technical Innovation, Strengths, Weaknesses

Method/Paper	Category	Key Technical Innovation	Strengths	Weaknesses	Benchmark/Task	Best Reported Performance	Year
Quantum Kernel SVM (General)	I. Feature Space Expansion	Mapping to representationally large $\mathcal{H}$, Q	High theoretical expressivity; established for classification	Highly sensitive to measurement noise; classical simulation bounds are often unclear [12].	General Classification	Theoretically Super-polynomial Speedup	2022
Variational Quantum Classifier (VQC)	IV. Parameterized Modeling	Hybrid optimization loop; parameterized unitary $U(\theta)$	Most hardware-accessible framework; adaptable to various tasks.	Prone to Barren Plateaus; performance degrades rapidly without physical constraints [5].	Iris Data Set	Feasible Proof-of-Concept	2022
Hamiltonian Simulation	II. Physics-Informed Processing	Quantum evolution operator e^{-iHt}	Advantage derived from fundamental physics; most robust theoretical grounding.	Requires high coherence and fault tolerance; complex initial state preparation.	Quantum Chemistry	Potential Exponential Speedup	2022
HHL	III. Distribution Learning	Quantum Fourier Transform; quantum point $P(\mathbf{x})$	Provides exponential speedup for sparse, well-conditioned data manifolds.	Data loading bottleneck; high-fidelity state preparation; sampling efficiency is low [18].	Linear Algebra Problems	Exponential Speedup	2009
Quantum-Enhanced SVM	I. Feature Space Expansion	Incorporating quantum state propagation into linear solvers	Provides a clear theoretical path for enhancing classical methods.	Often restricted to simplified, low-depth quantum components; advantage marginal on simple data [31].	Sparse Solutions	Minor Speedup	2020
PIML for PDE Solving	II. Physics-Informed Processing	Integrating PDE residuals into the loss function	Grounding ML in physical reality provides strong regularization.	Requires expert domain knowledge for correct constraint formulation.	Scientific Computing	Constrained Convergence	2021
Amplitude Encoding	IV. Parameterized Modeling	Mapping $\mathbf{x}$ to state amplitudes $\mathbb{S}$	$c(\mathbf{x})$	Efficient use of qubits; scales well with data dimension d	Data Encoding Theory	Theoretical Efficiency Gain	
Quantum CNNs (QCNNs)	IV. Parameterized Modeling	Enforcing local connectivity on the quantum feature space	Provides structure to the VQC, mitigating generic barren plateau issues.	The structure must be carefully chosen to ensure the desired inductive bias is maintained [12].	High Energy Physics Data	Improved Training Stability	2022
Quantum Generative Models (General)	III. Distribution Learning	Learning underlying correlation structures via quantum circuits	Potential for modeling complex physical correlation functions.	Lacks standardized metrics for measuring the "quantumness" of the learned distribution.	Data Synthesis	Qualitative Potential	2022
Quantum Fourier Map (General)	I. Feature Space Expansion	$K(\mathbf{x}, \mathbf{y})$ calculation via random measurements	Offers a path to explore high-dimensional feature spaces without explicit computation.	Measurement overhead and variance accumulation limit rigorous proof of advantage [12].	General ML Benchmarks	Requires Large Sample Size	2022
Variational Quantum Eigensolver (VQE)	IV. Parameterized Modeling	Hybrid optimization loop; parameterized unitary $U(\theta)$	Most hardware-accessible framework; adaptable to various tasks.	Prone to Barren Plateaus; performance degrades rapidly without physical constraints [5].	Molecular Energy	Chemical Accuracy	2022
Quantum Reservoir Computing	IV. Parameterized Modeling	Utilizing Hilbert space dynamics as memory/state evolution	Potential for complex time-series modeling using quantum memory effects.	The literature remains nascent; formal proofs of exponential advantage are rare [19].	Time Series Forecasting	Limited Empirical Evidence	2022
Quantum-Inspired Optimization	IV. Parameterized Modeling	Utilizing quantum concepts (e.g., quantum state fidelity) in classical optimization solvers	Bridges the gap between quantum theory and classical computational tools.	The advantage is often inspired rather than derived from true quantum computation.	Combinatorial Optimization	Effective Heuristic Improvement	2014
Quantum Differential Analysis	II. Physics-Informed Processing	Applying quantum principles to solve differential equations	Addresses foundational needs in theoretical physics and engineering.	Requires precise mapping of continuous differential operators to discrete quantum gates [18].	PDE Solving	Theoretical Pathway Established	2022
Quantum Image Recognition	IV. Parameterized Modeling	Utilizing QCNNs or VQC on preprocessed images	Demonstrates direct application to structured data types.	The performance gap relative to classical CNNs remains the central point of contention	Image Classification	Empirical Ambiguity	2021-2023
Quantum Optimization (General)	IV. Parameterized Modeling	Using QAOA/QA for combinatorial search	Direct application to NP-hard problems; clear theoretical target.	Mapping real-world constraints to the Ising model is non-trivial; scaling remains a major bottleneck [18].	Graph Coloring, Optimization	Theoretical Scaling Analysis	2022

6.1 Cross-Cutting Analysis: Overarching Trends and Insights

The comparison across these disparate methods reveals several critical, overarching trends that define the current state and trajectory of the field.

Firstly, there is a clear and persistent **Shift from Data-Centric to Physics-Centric Advantage**. Early enthusiasm in QML heavily focused on applying quantum circuits to general, classical datasets for classification, such as using QCNNs on image datasets [12] or VQCs on standard benchmarks like Iris [27]. However, the literature increasingly pivots toward problems where the difficulty is *inherent* to the system being modeled, such as molecular structure or plasma dynamics , signaling a maturation of the field's focus from algorithmic novelty to fundamental computational necessity.

Secondly, the **Optimization Bottleneck Outpaces Feature Space Complexity**. Despite the theoretical promise of exponentially large feature spaces offered by quantum kernels [16], the practical execution is consistently impeded by the optimization landscape. The sheer difficulty of calculating gradients, leading to phenomena like the Barren Plateau [5], suggests that the most significant immediate technological hurdle is not the encoding of information, but the *training* of the circuit that processes it. This has spurred research into structured ansätze [12] and constrained optimization methods [34].

A third crucial trend is the **Convergence on Hybridization as the Operational Standard**. Virtually all successful proposals—from QGANs analyzing distribution learning [37] to QCNNs for physics data [8]—rely on a classical optimizer managing a quantum evaluation subroutine. This confirms that the near-term quantum advantage will manifest not as a replacement for classical computing, but as a highly specialized, non-linear co-processor, where the quantum chip provides an intractable evaluation step that classical hardware cannot replicate efficiently.

Furthermore, a striking trend concerns the **Necessity of Quantum-Native Data**. The papers that exhibit the most robust theoretical arguments for advantage are those dealing with data whose structure is defined by quantum mechanics (e.g., electronic structure calculations, quantum field theories) [9]. Conversely, attempts to force quantum advantage onto classical data proxies often result in performance metrics that are comparable to, or worse than, state-of-the-art classical algorithms, suggesting that the computational gap $\mathrm{BQP} \setminus \mathrm{BPP}$ is only readily accessible through physical realism.

6.2 Core Trade-Offs in Quantum Machine Learning

The field is currently navigating several fundamental trade-offs that dictate the feasibility of any given quantum ML approach. The most prominent is the **Expressivity vs. Trainability Trade-off**. A circuit that possesses maximum theoretical expressivity, such as a highly deep or random quantum circuit, maximizes the model's capacity to fit any function (Universal Approximation Property, UAP) [12]. However, as established, this maximal expressivity often correlates with minimal gradient signal, rendering the circuit untrainable due to the Barren Plateau effect [5]. Conversely, circuits designed to be shallow and trainable (e.g., those incorporating strong, physical inductive biases) sacrifice some theoretical expressive power but gain immense practical stability. The literature suggests that for near-term success, constrained, physically-motivated designs are superior to maximally expressive, deep, but poorly behaved architectures.

Another vital trade-off is **Computational Resource Scaling vs. Implementation Feasibility**. In the context of quantum kernels, the potential advantage is tied to the exponential scaling of the feature space dimension. While this is theoretically powerful, realizing it requires the measurement of expectation values on exponentially large Hilbert spaces, which demands fault-tolerant quantum computation far beyond current NISQ capabilities. The current reality forces researchers to use randomized measurements [13], which effectively reduces the *proven* advantage to one that can be estimated statistically, creating a gap between the theoretical limit and the practical measurement ceiling.

Finally, there is the **General Parallelism vs. Structured Sequential Processing Trade-off**. Some approaches, like quantum simulation, inherently model complex, entangled, and highly parallel physical processes, suggesting massive inherent parallelism. In contrast, VQA frameworks often decompose these problems into sequential, parameter-by-parameter optimization steps. The literature suggests that while the underlying physics is massively parallel, the current *computational formalism* necessary to train the model forces a sequential, iterative optimization process, thereby masking the full potential parallelism available to the quantum hardware.

7 Open Challenges and Future Directions

The literature review confirms that while the theoretical potential is vast, the realization of robust, demonstrable quantum advantage requires breakthroughs in several deeply interconnected areas. These challenges demand coordinated efforts spanning theoretical physics, algorithm design, and quantum engineering.

7.1 Challenge 1: Developing Quantum-Native, Scalable Benchmarking Datasets

The single most significant impediment to achieving consensus on advantage is the lack of standardized, large-scale datasets derived from quantum phenomena. Current reliance on classical proxies—such as image datasets or simple numerical sequences—provides an insufficient test bed [6]. To move forward, the field requires benchmark problems where the computational hardness *cannot* be mapped onto a classical polynomial-time algorithm, and where the complexity is demonstrably linked to the structure of quantum physics. A promising direction involves leveraging the Hamiltonian simulation paradigm. Instead of using abstract data, future benchmarks should involve simulating increasingly complex physical systems, such as multi-electron molecular interactions or lattice gauge theories, where the ground-state energy calculation itself serves as the objective function . Furthermore, incorporating ideas from information theory, as highlighted by necessary tools for verifying advantage [22], must become part of the benchmark specification.

7.2 Challenge 2: Developing Theory for Quantum Advantage Quantification

The current framework lacks a general, complexity-theoretic criterion for advantage. Researchers are currently forced to appeal to ad-hoc arguments based on feature space size or Hamiltonian complexity. Future work must establish a formal mathematical framework that precisely delineates the boundary between BQP and classical classes, allowing for a definitive 'quantum advantage proof' for a given problem class. This requires moving beyond demonstrating *potential* speedup to proving *necessary* speedup. One actionable direction is to formalize the "quantum cost function," quantifying the minimal resource (qubits, gates, shots) required for a quantum computation versus the best

known classical simulation, thereby creating a quantifiable metric that guides algorithm design, as proposed conceptually by analyzing information-theoretic bounds [14].

7.3 Challenge 3: Designing Truly Robust and Scalable Quantum Compilers/Middleware

The current research often treats quantum algorithms and ML optimization as separate components. The true breakthrough requires the development of a unified, intelligent software stack—the "Quantum Advantage Seeker." This middleware must ingest a high-level scientific problem (e.g., "Determine the optimal catalyst structure") and automatically select the optimal computational pathway. This selection must dynamically weigh options: Is the problem best solved by a quantum kernel estimation [16]? Should it be constrained by a known PDE via a PIML approach ? Or is a parameterized circuit the only viable option given current hardware constraints [4]? This mandates research in automated algorithm synthesis, potentially drawing inspiration from symbolic AI techniques and advanced resource management, moving beyond simple quantum circuit compilation [1].

7.4 Challenge 4: Mitigating Optimization Inherent Failures in Variational Models

The persistent problem of Barren Plateaus represents a fundamental limitation in the current operational VQA paradigm. Overcoming this requires moving beyond merely structuring the circuit (e.g., QCNNs [12]) and developing novel optimization theories. Two specific directions are critical: first, developing gradient estimators that are provably robust to the circuit depth, perhaps by localizing the gradient calculation to specific physical observables [18]. Second, integrating knowledge from non-classical optimization fields, such as those studied in quantum computing for optimization [30], into the ML loop to guide the search away from flat regions of the cost landscape. The goal is to achieve *trainability* parity with the theoretical *expressivity*.

7.5 Challenge 5: Bridging Quantum Computation with Classical Data Structures

While the optimal path favors quantum-native data, the immediate commercial and research utility demands handling large volumes of classical data. The challenge here is to develop rigorous, measurable mechanisms that guarantee a quantum advantage when the input data ($\mathbf{x}$) is classical and arbitrary. Future work must focus on developing quantum data embeddings that are provably more powerful than classical feature mapping techniques, perhaps by exploiting entanglement structures in the embedding process itself, similar to how specialized quantum state preparation protocols are beginning to be explored [31]. This involves quantifying the information gain provided by the quantum state preparation relative to the classical information content of the input vector.

8 Limitations of This Survey

This survey provides a high-level taxonomy of the field rather than an exhaustive, empirical results database. Consequently, it necessarily focuses on *theoretical* and *potential* advantages, meaning that many comparisons are based on idealized model assumptions rather than direct, end-to-end benchmarks on state-of-the-art quantum hardware. The literature review is necessarily constrained by the published corpus, leading to an inherent bias toward methods that are mathematically tractable or have been recently published in high-impact venues. Furthermore, the distinction between quantum advantage derived from *algorithmic complexity* versus advantage derived from *quantum physical simulation* remains a zone of active debate, and the survey synthesizes the current leading theories without resolving this meta-theoretical conflict.

9 Conclusion

This survey has successfully constructed a taxonomy for Quantum Machine Learning by shifting the focus from the algorithm employed to the mathematical source of the required computational hardness: be it the exponential dimensionality of a feature space, the underlying Hamiltonian of a physical system, the complexity of a joint probability distribution, or the geometric challenges of optimization. The analysis unambiguously demonstrates that the most compelling and theoretically robust pathway toward exponential quantum advantage lies at the intersection of physics simulation and machine learning, where physical laws provide the necessary constraints to stabilize the learning process. While the field has achieved impressive milestones in demonstrating the *feasibility* of hybrid

quantum-classical training loops, the consensus points to the fact that the field has yet to achieve the necessary standardized, quantum-native benchmarking datasets or the generalized complexity theory required to move from hypothesis to proven capability. The trajectory clearly indicates that the next breakthrough will not stem from an improved VQA ansatz alone, but from the creation of a sophisticated, multi-modal middleware capable of diagnosing and selecting the mathematically optimal, physics-guided quantum computational strategy for a given scientific problem.

References

[1] Israel F. Araujo, Daniel K. Park, Francesco Petruccione, and Adenilton J. da Silva. A divide-and-conquer algorithm for quantum state preparation. *Scientific Reports*, 2021. doi: 10.1038/s41598-021-85474-1. URL https://doi.org/10.1038/s41598-021-85474-1.

[2] Kishor Bharti, Alba Cervera-Lierta, Thi Ha Kyaw, Tobias Haug, Sumner Alperin-Lea, Abhinav Anand, Matthias Degroote, Hermanni Heimonen, Jakob S. Kottmann, Tim Menke, Wai-Keong Mok, Sukin Sim, L. C. Kwek, and Alán Aspuru-Guzik. Noisy intermediate-scale quantum (nisq) algorithms. *arXiv (Cornell University)*, 2021. doi: 10.48550/arxiv.2101.08448. URL https://doi.org/10.48550/arxiv.2101.08448.

[3] Kishor Bharti, Alba Cervera-Lierta, Thi Ha Kyaw, Tobias Haug, Sumner Alperin-Lea, Abhinav Anand, Matthias Degroote, Hermanni Heimonen, Jakob S. Kottmann, Tim Menke, Wai-Keong Mok, Sukin Sim, L. C. Kwek, and Alán Aspuru-Guzik. Noisy intermediate-scale quantum algorithms. *Reviews of Modern Physics*, 2022. doi: 10.1103/revmodphys.94.015004. URL https://doi.org/10.1103/revmodphys.94.015004.

[4] M. Cerezo, Andrew Arrasmith, Ryan Babbush, Simon C. Benjamin, Suguru Endo, Keisuke Fujii, Jarrod R. McClean, Kosuke Mitarai, Xiao Yuan, Lukasz Cincio, and Patrick J. Coles. Variational quantum algorithms. *Nature Reviews Physics*, 2021. doi: 10.1038/s42254-021-00348-9. URL https://doi.org/10.1038/s42254-021-00348-9.

[5] M. Cerezo, Akira Sone, Tyler Volkoff, Łukasz Cincio, and Patrick J. Coles. Cost function dependent barren plateaus in shallow parametrized quantum circuits. *Nature Communications*, 2021. doi: 10.1038/s41467-021-21728-w. URL https://doi.org/10.1038/s41467-021-21728-w.

[6] M. Cerezo, Guillaume Verdon, Hsin-Yuan Huang, Łukasz Cincio, and Patrick J. Coles. Challenges and opportunities in quantum machine learning. *Nature Computational Science*, 2022. doi: 10.1038/s43588-022-00311-3. URL https://doi.org/10.1038/s43588-022-00311-3.

[7] Simantini Chakraborty, Tamal Das, Saurav Sutradhar, Mrinmoy Das, and Suman Deb. An analytical review of quantum neural network models and relevant research. 2020. doi: 10.1109/icces48766.2020.9137960. URL https://doi.org/10.1109/icces48766.2020.9137960.

[8] Samuel Yen-Chi Chen, Tzu-Chieh Wei, C. Zhang, Haiwang Yu, and Shinjae Yoo. Quantum convolutional neural networks for high energy physics data analysis. *Physical Review Research*, 2022. doi: 10.1103/physrevresearch.4.013231. URL https://doi.org/10.1103/physrevresearch.4.013231.

[9] Andrew J. Daley, Immanuel Bloch, Christian Kokail, Stuart Flannigan, Natalie Pearson, Matthias Troyer, and P. Zoller. Practical quantum advantage in quantum simulation. *Nature*, 2022. doi: 10.1038/s41586-022-04940-6. URL https://doi.org/10.1038/s41586-022-04940-6.

[10] Daniel J. Egger, Claudio Gambella, Jakub Marecek, Scott McFaddin, Martin Mevissen, Rudy Raymond, Andrea Simonetto, Stefan Woerner, and Elena Yndurain. Quantum computing for finance: State-of-the-art and future prospects. *IEEE Transactions on Quantum Engineering*, 2020. doi: 10.1109/tqe.2020.3030314. URL https://doi.org/10.1109/tqe.2020.3030314.

[11] Dar Gilboa, Jarrod R. McClean, Daniel Soudry, and Jarrod R. McClean. Exponential quantum communication advantage in distributed inference and learning. *arXiv (Cornell University)*, 2023. doi: 10.48550/arxiv.2310.07136. URL https://doi.org/10.48550/arxiv.2310.07136.

[12] Takahiro Goto, Quoc Hoan Tran, and Kohei Nakajima. Universal approximation property of quantum machine learning models in quantum-enhanced feature spaces. *Physical Review Letters*, 2021. doi: 10.1103/physrevlett.127.090506. URL https://doi.org/10.1103/physrevlett.127.090506.

[13] Tobias Haug, Chris N. Self, and M. S. Kim. Quantum machine learning of large datasets using randomized measurements. *Machine Learning Science and Technology*, 2023. doi: 10.1088/2632-2153/acb0b4. URL https://doi.org/10.1088/2632-2153/acb0b4.

[14] Hsin-Yuan Huang, Richard Kueng, and John Preskill. Information-theoretic bounds on quantum advantage in machine learning. *Physical Review Letters*, 2021. doi: 10.1103/physrevlett.126.190505. URL https://doi.org/10.1103/physrevlett.126.190505.

[15] Kaixuan Huang, Zheng-An Wang, Chao Song, Kai Xu, Hekang Li, Zhen Wang, Qiujiang Guo, Zixuan Song, Zhibo Liu, Dongning Zheng, Dong-Ling Deng, H. Wang, Jianguo Tian, and Heng Fan. Quantum generative adversarial networks with multiple superconducting qubits. *npj Quantum Information*, 2021. doi: 10.1038/s41534-021-00503-1. URL https://doi.org/10.1038/s41534-021-00503-1.

[16] Jonas Jäger and Roman V. Krems. Universal expressiveness of variational quantum classifiers and quantum kernels for support vector machines. *Nature Communications*, 2023. doi: 10.1038/s41467-023-36144-5. URL https://doi.org/10.1038/s41467-023-36144-5.

[17] William D. Kalfus, Guilhem Ribeill, Graham E. Rowlands, Hari Krovi, Thomas Ohki, and Luke C. G. Govia. Hilbert space as a computational resource in reservoir computing. *Physical Review Research*, 2022. doi: 10.1103/physrevresearch.4.033007. URL https://doi.org/10.1103/physrevresearch.4.033007.

[18] Bobak T. Kiani, Giacomo De Palma, Milad Marvian, Zi-Wen Liu, and Seth Lloyd. Learning quantum data with the quantum earth mover's distance. *Quantum Science and Technology*, 2022. doi: 10.1088/2058-9565/ac79c9. URL https://doi.org/10.1088/2058-9565/ac79c9.

[19] Florian J. Kiwit, Maximilian A. Wolf, Marwa Marso, Philipp Ross, J. Lorenz, Carlos A. Riofrío, and André Luckow. Benchmarking quantum generative learning: A study on scalability and noise resilience using quark. *KI - Künstliche Intelligenz*, 2024. doi: 10.1007/s13218-024-00864-7. URL https://doi.org/10.1007/s13218-024-00864-7.

[20] Tong Li, Shibin Zhang, and Jinyue Xia. Quantum generative adversarial network: A survey. *Computers, materials & continua/Computers, materials & continua (Print)*, 2020. doi: 10.32604/cmc.2020.010551. URL https://doi.org/10.32604/cmc.2020.010551.

[21] Jie Lin, Dan-Bo Zhang, Shuo Zhang, Tan Li, Xiang Wang, and Wan-Su Bao. Quantum-enhanced least-square support vector machine: Simplified quantum algorithm and sparse solutions. *Physics Letters A*, 2020. doi: 10.1016/j.physleta.2020.126590. URL https://doi.org/10.1016/j.physleta.2020.126590.

[22] Yong Liu, Yaojian Chen, Chu Guo, Jiawei Song, Xinmin Shi, Lin Gan, Wenzhao Wu, Wei Wu, Haohuan Fu, Xin Liu, Dexun Chen, Zhifeng Zhao, Guangwen Yang, and Jiangang Gao. Verifying quantum advantage experiments with multiple amplitude tensor network contraction. *Physical Review Letters*, 2024. doi: 10.1103/physrevlett.132.030601. URL https://doi.org/10.1103/physrevlett.132.030601.

[23] Yunchao Liu, Srinivasan Arunachalam, and Kristan Temme. A rigorous and robust quantum speed-up in supervised machine learning. *Nature Physics*, 2021. doi: 10.1038/s41567-021-01287-z. URL https://doi.org/10.1038/s41567-021-01287-z.

[24] Francesco Di Marcantonio, Massimiliano Incudini, Davide Tezza, and Michele Grossi. Quantum advantage seeker with kernels (quask): a software framework to speed up the research in quantum

machine learning. *Quantum Machine Intelligence*, 2023. doi: 10.1007/s42484-023-00107-2. URL https://doi.org/10.1007/s42484-023-00107-2.

[25] Stefano Mensa, Emre Sahin, Francesco Tacchino, Panagiotis Kl. Barkoutsos, and Ivano Tavernelli. Quantum machine learning framework for virtual screening in drug discovery: a prospective quantum advantage. *Machine Learning Science and Technology*, 2023. doi: 10.1088/2632-2153/acb900. URL https://doi.org/10.1088/2632-2153/acb900.

[26] Yuta Mizuno and Tamiki Komatsuzaki. Finding optimal pathways in chemical reaction networks using ising machines. *Physical Review Research*, 2024. doi: 10.1103/physrevresearch.6.013115. URL https://doi.org/10.1103/physrevresearch.6.013115.

[27] Ilya Piatrenka and Marian Rusek. Quantum variational multi-class classifier for the iris data set. *Lecture notes in computer science*, 2022. doi: 10.1007/978-3-031-08760-8_21. URL https://doi.org/10.1007/978-3-031-08760-8_21.

[28] Mohan Raja Pulicharla. Hybrid quantum-classical machine learning models: Powering the future of ai. *Journal of Science & Technology*, 2023. doi: 10.55662/jst.2023.4102. URL https://doi.org/10.55662/jst.2023.4102.

[29] Sebastian Raubitzek and Kevin Mallinger. On the applicability of quantum machine learning. *Entropy*, 2023. doi: 10.3390/e25070992. URL https://doi.org/10.3390/e25070992.

[30] Krishanu Sankar, Artur Scherer, Satoshi Kako, Sam Reifenstein, Navid Ghadermarzy, Willem B. Krayenhoff, Yoshitaka Inui, Edwin Ng, Tatsuhiro Onodera, Pooya Ronagh, and Y. Yamamoto. A benchmarking study of quantum algorithms for combinatorial optimization. *npj Quantum Information*, 2024. doi: 10.1038/s41534-024-00856-3. URL https://doi.org/10.1038/s41534-024-00856-3.

[31] Alok Shukla and Prakash Vedula. An efficient quantum algorithm for preparation of uniform quantum superposition states. In *Quantum Information Processing*, 2024. doi: 10.1007/s11128-024-04258-4. URL https://doi.org/10.1007/s11128-024-04258-4.

[32] Lucas Slattery, Ruslan Shaydulin, Shouvanik Chakrabarti, Marco Pistoia, Sami Khairy, and Stefan M. Wild. Numerical evidence against advantage with quantum fidelity kernels on classical data. *Physical review. A/Physical review, A*, 2023. doi: 10.1103/physreva.107.062417. URL https://doi.org/10.1103/physreva.107.062417.

[33] Supanut Thanasilp, Samson Wang, M. Cerezo, and Zoë Holmes. Exponential concentration in quantum kernel methods. *Nature Communications*, 2024. doi: 10.1038/s41467-024-49287-w. URL https://doi.org/10.1038/s41467-024-49287-w.

[34] Yabo Wang, Bo Qi, Christopher Ferrie, and Daoyi Dong. Trainability enhancement of parameterized quantum circuits via reduced-domain parameter initialization. *Physical Review Applied*, 2024. doi: 10.1103/physrevapplied.22.054005. URL https://doi.org/10.1103/physrevapplied.22.054005.

[35] Amine Zeguendry, Zahi Jarir, and Mohamed Quafafou. Quantum machine learning: A review and case studies. *Entropy*, 2023. doi: 10.3390/e25020287. URL https://doi.org/10.3390/e25020287.

[36] Nanrun Zhou, Tianfeng Zhang, Xinwen Xie, and Junyun Wu. Hybrid quantum–classical generative adversarial networks for image generation via learning discrete distribution. In *Signal Processing Image Communication*, 2022. doi: 10.1016/j.image.2022.116891. URL https://doi.org/10.1016/j.image.2022.116891.

[37] Elton Yechao Zhu, Sonika Johri, Dave Bacon, Mert Esencan, Jungsang Kim, Mark Muir, Nikhil Murgai, Jason Nguyen, Neal C. Pisenti, Adam Schouela, Ksenia Sosnova, and Ken Wright. Generative quantum learning of joint probability distribution functions. *Physical*

Review Research, 2022. doi: 10.1103/physrevresearch.4.043092. URL https://doi.org/10.1103/physrevresearch.4.043092.