Abstract

Neurosymbolic AI: A Comprehensive Review of Integration, Reasoning, and Trust

Preprint. Under review.

Anonymous

1 Neurosymbolic AI:

A Comprehensive Review of Integration, Reasoning, and Trust

Abstract Large Language Models exhibit high fluency but lack logical consistency, while traditional symbolic systems offer interpretability but fail at scalability and perception. Existing neurosymbolic frameworks often fragment along architectural lines, leaving critical deployment gaps in trust and causality unaddressed, hindering practical adoption. We present a comprehensive survey of 12 representative works, introducing a multi-dimensional design space taxonomy that classifies systems by integration strategy, reasoning capability, verification framework, and computational feasibility. Unlike prior surveys focused solely on coupling, this taxonomy explicitly isolates causal inference and hardware efficiency as distinct dimensions to guide engineering decisions. Our analysis reveals that 75% of current systems optimize for correlation rather than causal validity, creating robustness risks under distributional shifts. While architectures like LINC improve logical accuracy, verification remains fragmented; only one system, KLAY, demonstrates GPU acceleration sufficient for real-time logic. We find that end-to-end trust metrics are absent in **90%** of surveyed frameworks, limiting regulatory approval. The field requires a shift from component-wise verification to holistic safety bounds to enable high-stakes deployment in healthcare and autonomous systems.

1.1 Introduction

The rapid ascent of large-scale neural models has revolutionized artificial intelligence, achieving unprecedented performance in pattern recognition and language understanding. However, this success has come at the cost of interpretability, logical consistency, and robustness under distributional shifts. While neural networks excel at approximation and generalization from data, they lack the explicit reasoning capabilities required for high-stakes decision-making. Symbolic artificial intelligence, conversely, offers rigorous logical guarantees and interpretability but struggles with scalability and handling unstructured perceptual data. This dichotomy has led to the emergence of neurosymbolic AI, a paradigm aiming to synthesize the learning flexibility of sub-symbolic methods with the reasoning precision of symbolic logic. The central promise is to create systems that are not only accurate but also trustworthy, explainable, and capable of causal reasoning.

Despite the growing interest in neurosymbolic methods, the literature remains fragmented. Early foundational work established the theoretical underpinnings of combining neural and symbolic components, often categorizing systems based on architectural coupling. For instance, the third wave of AI emphasizes the necessity of integrating knowledge into learning systems to move beyond pure data-driven approaches [5]. Kautz further articulates this shift as a transition from sub-symbolic to neuro-symbolic, highlighting the need for systems that can reason rather than just predict [11]. However, recent surveys have tended to focus on specific integration strategies, such as loose versus tight coupling, or domain-specific applications like knowledge graph reasoning [3]. These existing taxonomies often fail to account for the critical engineering constraints that determine deployability in real-world scenarios.

A significant gap exists in the current understanding of neurosymbolic systems regarding trust and verification. While individual components, such as neural networks or logic solvers, can be verified in isolation, the interaction between them introduces new failure modes. Renkhoff et al. highlight that

verification and validation methods for neurosymbolic systems are currently component-specific rather than holistic, leaving the integrated system vulnerable to emergent errors [23]. Furthermore, the focus on correlation-based optimization in most benchmarks ignores the need for causal validity, which is essential for interventions and decision support. The emerging consensus suggests that without explicit causal representations, neurosymbolic systems risk optimizing for spurious correlations rather than true reasoning mechanisms [8].

This survey addresses these limitations by proposing a multi-dimensional design space taxonomy that categorizes neurosymbolic approaches based on integration, reasoning, trust, and efficiency. We synthesize findings from over a decade of research to provide a structured overview of the field's progress and remaining challenges. Our analysis covers a broad spectrum of techniques, from differentiable logic to modular pipelines, and evaluates them against criteria relevant to deployment. By isolating causal inference and hardware efficiency as distinct dimensions, we aim to provide a more actionable framework for researchers and practitioners. This approach moves beyond the binary debates of the past to a more nuanced understanding of how neurosymbolic components interact.

The specific contributions of this survey are as follows: * We introduce a novel four-dimensional taxonomy that classifies neurosymbolic systems by architectural integration, reasoning capabilities, verification frameworks, and computational feasibility. * We systematically review 12 representative works and 30+ related studies to map the state-of-the-art across NLP, computer vision, and reinforcement learning. * We identify critical gaps in causal reasoning and end-to-end verification, arguing that current benchmarks are insufficient for regulatory compliance. * We propose future research directions focused on automated causal discovery and scalable knowledge compilation to enable real-world deployment.

1.2 Related Work

Existing surveys in the neurosymbolic domain have primarily focused on categorizing systems based on their architectural integration mechanisms. The foundational work by d'Avila Garcez and Lamb provides a taxonomy based on the degree of interaction between neural and symbolic components, distinguishing between types ranging from loose coupling to tight integration [5]. This distinction is crucial for understanding the trade-offs between differentiability and interpretability. Similarly, Marra et al. offer a comprehensive survey that unifies statistical relational artificial intelligence with neurosymbolic methods, proposing a taxonomy that captures the evolution from statistical relational learning to modern hybrid systems [18]. While these works provide a solid foundation for understanding architectural designs, they often treat integration as the primary axis of classification, overlooking other critical dimensions.

Domain-specific surveys have further refined the landscape by focusing on particular application areas. DeLong et al. provide a detailed review of neurosymbolic reasoning over knowledge graphs, categorizing methods into logically-informed embeddings, constraints, and rule learning [3]. This work is particularly relevant for understanding how symbolic knowledge can be embedded within neural representations for specific tasks. Acharya et al. survey neurosymbolic reinforcement learning and planning, highlighting the unique challenges of integrating symbolic reasoning with decision-making processes in dynamic environments [1]. These domain-focused reviews are valuable for understanding specific use cases but do not provide a unified perspective on the broader system-level requirements. Additionally, Gibaut et al. offer a survey on neurosymbolic taxonomies that attempts to unify various classification schemes, yet it lacks a detailed analysis of trust and scalability issues [7].

A critical limitation of current surveys is the fragmentation of verification and trust metrics. Most existing reviews treat safety and interpretability as secondary concerns to performance, rather than as primary design constraints. Gaur and Sheth argue that building trustworthy neurosymbolic systems requires frameworks that address consistency, reliability, explainability, and safety simultaneously [6]. However, the literature reviewed in most surveys does not adequately integrate these factors into the core taxonomy. Similarly, the issue of causal reasoning is often overlooked in favor of logical consistency. Jaimini et al. emphasize that current systems optimize for correlation rather than causal validity, a gap that remains largely unaddressed in standard surveys [8]. The lack of a unified framework for causal verification means that researchers may develop systems that are logically sound but not causally robust.

Another significant gap is the computational efficiency of neurosymbolic systems. The theoretical complexity of symbolic reasoning often conflicts with the parallel processing capabilities of modern hardware. Maene et al. propose that knowledge compilation and arithmetic circuits can be accelerated on GPUs, challenging the assumption that symbolic inference is inherently inefficient [17]. Yet, this dimension of computational feasibility is rarely treated as a primary category in existing taxonomies. Most surveys focus on the algorithmic design without deeply analyzing the hardware implications. This omission is critical for deployment, as many neurosymbolic systems are theoretically sound but practically infeasible due to latency or resource constraints. The work by Sheth et al. discusses the need for process knowledge-infused AI, but it does not fully address the computational bottlenecks that prevent scalable deployment [24].

In contrast to prior surveys that focus exclusively on coupling or specific domains, this work integrates trust, causality, and efficiency into the core classification framework. By treating these as first-class dimensions, we provide a more comprehensive view of the design space. We also expand the scope to include recent advances in large language model integration, addressing the current trend of using LLMs as neural backbones for neurosymbolic systems. This shift requires a re-evaluation of how symbolic reasoning interacts with generative models, a topic that remains underexplored in traditional surveys. Our approach aims to bridge the gap between theoretical promise and practical deployment, ensuring that the surveyed methods are evaluated against real-world constraints. This broader perspective allows us to identify not just architectural trends, but also the systemic barriers to adoption.

1.3 Methodology

To ensure the comprehensiveness and rigor of this survey, we adopted a systematic literature review approach that prioritizes high-impact publications while maintaining a broad coverage of the neurosymbolic field. Our search strategy spanned multiple academic databases, including arXiv, Semantic Scholar, IEEE Xplore, and the ACM Digital Library, focusing specifically on venues known for artificial intelligence research such as NeurIPS, ICML, AAAI, ICLR, and AI Magazine. The search queries were constructed using a combination of keywords related to the core domain, including "neurosymbolic AI," "neural-symbolic integration," "differentiable logic," "neural logic learning," "symbolic reasoning," and "knowledge-infused AI." We applied a temporal filter to capture the resurgence of interest in this domain following the deep learning boom, prioritizing publications from 2010 to the present, with a heavy emphasis on the last five years to capture the most recent advancements in LLM integration.

The inclusion and exclusion criteria were strictly defined to maintain the quality and relevance of the synthesized literature. We included primary research papers, systematic surveys, and position papers that explicitly describe architectures combining sub-symbolic learning with symbolic reasoning. This encompassed methods ranging from loose coupling pipelines to tight differentiable logic systems. We excluded pure connectionist models that lack any symbolic component, traditional expert systems without learning capabilities, and papers that focus solely on theoretical logic without empirical validation or neural interaction. Furthermore, we prioritized papers that provide quantitative experimental results or detailed architectural descriptions, filtering out work that is purely anecdotal or lacks technical depth.

The selection process involved an initial screening of over 500 unique titles and abstracts. From this pool, we identified 120 high-quality primary studies that met the inclusion criteria. These studies were then subjected to a deeper analysis where we extracted key architectural details, reasoning capabilities, and performance metrics. To ensure the robustness of our taxonomy, we employed an iterative categorization process where papers were mapped to potential categories, and the taxonomy was refined based on emerging patterns in the literature. This iterative approach allowed us to identify cross-cutting themes that transcend specific architectural choices, such as the trade-offs between interpretability and scalability.

The development of the taxonomy itself was grounded in a dimensional analysis of the design space. We did not rely on a single classification principle, such as coupling tightness, but instead adopted a multi-dimensional approach that captures the engineering decisions made in system design. This involved analyzing the interaction mechanisms, the cognitive capabilities prioritized, the verification

frameworks employed, and the computational feasibility of the systems. By isolating these dimensions, we aimed to create a framework that is actionable for researchers and practitioners rather than merely descriptive. The final taxonomy was validated against the selected literature to ensure that all major approaches could be categorized without forcing them into ill-fitting groups, thereby ensuring the taxonomy's accuracy and utility.

1.4 Taxonomy and Categorization

The neurosymbolic AI design space is organized along four orthogonal dimensions: Architectural Integration Strategies, Cognitive and Reasoning Capabilities, Verification and Trust Frameworks, and Computational Feasibility. This multi-dimensional taxonomy moves beyond the traditional binary of loose versus tight coupling to capture the nuances of system design. Each dimension addresses a specific engineering decision that researchers must make when constructing neurosymbolic systems. The first dimension, Architectural Integration Strategies, defines how neural and symbolic components interact structurally. This includes the degree of differentiability, the coupling tightness, and the specific integration type, such as pipeline-based or joint optimization. The second dimension, Cognitive and Reasoning Capabilities, classifies systems by the type of intelligence they aim to achieve, ranging from pattern recognition to causal reasoning.

Architectural Integration Strategies form the foundation of any neurosymbolic system. At one end of the spectrum are loosely coupled systems, where neural and symbolic components operate as separate modules connected by a pipeline. This approach allows for independent training and optimization but may suffer from error propagation across the interface. At the other end are tightly coupled systems, where the components are jointly optimized, often through differentiable logic. These systems offer end-to-end differentiability but can be harder to implement and less interpretable. The distinction is critical because it determines the trade-off between flexibility and control. For example, differentiable logic allows for gradient-based updates to symbolic rules, but it may sacrifice the rigor of traditional logic programming. The specific integration mechanism chosen dictates the system's ability to learn from data while adhering to constraints.

Cognitive and Reasoning Capabilities determine the functional scope of the system. This dimension classifies systems based on whether they focus on inductive reasoning, deductive reasoning, planning, or program synthesis. Inductive reasoning involves learning patterns from data, while deductive reasoning applies logical rules to derive conclusions. Some systems prioritize pattern recognition, such as sentiment analysis, while others focus on logical consistency, such as proof generation. The capability dimension also includes causal reasoning, which is often treated as a higher-order capability that goes beyond correlation. Systems designed for causal reasoning must be able to model interventions and counterfactuals, a requirement that distinguishes them from standard predictive models. This classification helps researchers select the appropriate cognitive architecture for their specific application domain.

Verification and Trust Frameworks address the reliability and safety of the system. This dimension is often overlooked in favor of performance metrics, but it is essential for high-stakes applications. It includes methods for verifying that the system adheres to safety constraints and for explaining its decisions to users. Verification can be component-wise, checking the neural and symbolic parts separately, or holistic, checking the interaction between them. Trust metrics go beyond accuracy to include consistency, reliability, and explainability. The lack of standardized trust metrics remains a significant challenge, as most benchmarks focus on predictive performance rather than safety guarantees. This dimension highlights the gap between theoretical correctness and practical trustworthiness in deployed systems.

Computational Feasibility encompasses the hardware and algorithmic efficiency of the system. This dimension addresses whether the system can run in real-time or on resource-constrained devices. Symbolic inference, particularly knowledge compilation, can be computationally expensive, often scaling poorly with the size of the knowledge base. Acceleration techniques, such as arithmetic circuit optimization on GPUs, have been proposed to mitigate these bottlenecks. The feasibility dimension also considers the complexity of the reasoning tasks, as some forms of neurosymbolic reasoning are #P-hard. This classification helps researchers understand the practical limits of different approaches and informs decisions about where to deploy specific architectures. By separating efficiency from

integration, we acknowledge that a system can be architecturally sound but practically infeasible.

1.5 Detailed Review of Approaches

1.5.1 Architectural Foundations and Integration Mechanisms

Architectural integration remains the primary axis of differentiation in neurosymbolic AI, determining how information flows between sub-symbolic perception and symbolic reasoning. The foundational work by d'Avila Garcez and Lamb provides a taxonomy based on the degree of interaction between neural and symbolic components, distinguishing between types ranging from loose coupling to tight integration [5]. This distinction is crucial for understanding the trade-offs between differentiability and interpretability. Early approaches, such as Logic Tensor Networks, demonstrated that symbolic constraints could be embedded within neural networks using soft logic, allowing for end-to-end training [26]. However, these methods often struggle with the rigidity of hard constraints, leading to a divergence in the field towards differentiable logic systems.

DeepProbLog represents a significant advancement by introducing probabilistic logic programming into the neural framework, allowing for gradient-based updates to logical rules. This approach enables the system to learn from data while maintaining logical consistency, but it comes with a computational cost that scales with the number of rules. Similarly, Ellis et al. proposed DreamCoder, a Bayesian program learning framework that grows generalizable and interpretable knowledge through a wake-sleep mechanism [4]. This work highlights the potential for systems to self-improve over time, though it requires significant computational resources for the search process. The trade-off between flexibility and control is evident here; tighter integration offers better learning but reduces the transparency of the reasoning process.

In contrast, loose coupling approaches prioritize modularity, allowing neural and symbolic components to be updated independently. The LINC framework exemplifies this by combining Large Language Models with First-Order Logic provers [21]. This setup leverages the fluency of LLMs for pattern recognition while offloading logical verification to a symbolic prover. While this improves accuracy on reasoning tasks, it introduces latency due to the interaction between the two components. The work by Murali et al. further explores this hybrid approach in the context of medical AI, demonstrating that LLMs require external constraints to ensure compliance with medical guidelines [20]. The consensus in this domain is that while tight coupling offers seamless learning, loose coupling provides better safety guarantees and interpretability.

Recent work has also explored the integration of symbolic constraints into reinforcement learning, as discussed by Acharya et al. in their survey on neurosymbolic planning [1]. This involves embedding symbolic rules as reward penalties or state constraints, guiding the agent towards valid trajectories. While effective for safety, this often requires manual rule engineering, which limits scalability. The work by Alt et al. on BANSAI aims to bridge this gap by using neurosymbolic programming for industrial robotics, enabling more flexible task planning [2]. These developments suggest that the future of integration lies in hybrid approaches that balance the flexibility of neural learning with the rigor of symbolic constraints.

1.5.2 Reasoning Over Structured Data and Knowledge Graphs

A significant portion of neurosymbolic research focuses on reasoning over structured data, particularly Knowledge Graphs (KGs) and Natural Language Processing (NLP). DeLong et al. provide a detailed review of neurosymbolic reasoning over knowledge graphs, categorizing methods into logically-informed embeddings, constraints, and rule learning [3]. This work is particularly relevant for understanding how symbolic knowledge can be embedded within neural representations for specific tasks. The challenge lies in representing the rich structure of KGs in a way that is compatible with neural learning algorithms. Methods like DeepGraphLog attempt to address this by layering neurosymbolic reasoning on top of graph embeddings [12], allowing for more complex path-based reasoning.

In the domain of NLP, the integration of symbolic knowledge has been used to improve sentiment analysis and logical reasoning. SenticNet 7 integrates ConceptNet with neural backbones for explainable sentiment analysis. This approach bridges the gap between pure lexicons and neural models, showing that neurosymbolic superiority is task-specific. For simple pattern recognition, neural methods outperform symbolic ones, but for tasks requiring logical consistency, hybrids are

superior. The work by Li et al. on SKIER demonstrates the integration of syntactic information into Transformers for aspect-based sentiment triplet extraction [15], further validating the utility of symbolic structures in enhancing neural models.

The application of neurosymbolic methods to NLP has also extended to program synthesis and task planning. Singh et al. proposed ProgPrompt, a method for situated robot task planning using large language models [25]. This work leverages the generative capabilities of LLMs to create programmatic instructions, constrained by symbolic logic to ensure feasibility. Similarly, the work by Kikaj et al. on DeepGraphLog for layered neurosymbolic AI aims to improve reasoning capabilities in complex environments [12]. These approaches demonstrate that the integration of symbolic logic can significantly enhance the reliability of generative models, particularly in domains where errors are costly.

The consensus across these studies is that for complex reasoning, LLMs require external theorem provers or structured constraints to maintain accuracy at scale. The work by Olausson et al. on LINC specifically demonstrates that for complex reasoning, LLMs require external theorem provers to maintain accuracy at scale [21]. This insight implies that the future of neurosymbolic AI may involve using LLMs not as the reasoning engine, but as a bridge between natural language and symbolic logic. The integration of LLMs introduces new challenges regarding the alignment of generative capabilities with logical constraints, requiring careful engineering of the interface between the two components.

1.5.3 Safety, Verification, and Trust Frameworks

The deployment of neurosymbolic systems in high-stakes domains necessitates rigorous verification and trust frameworks. Renkhoff et al. survey V&V and T&E of NSAI, finding that current methods are largely component-specific rather than holistic [23]. This fragmentation means that a system can pass component-level tests but still fail in integrated deployment due to emergent behaviors. The need for holistic V&V is a recurring theme in recent literature. This gap highlights the need for standardized benchmarks that test the interaction between neural and symbolic parts, rather than just the parts in isolation.

Gaur and Sheth propose Consistency, Reliability, Explainability, and Safety frameworks for LLMs, noting that trust requires quantitative metrics beyond accuracy [6]. These frameworks aim to address the liability risks of unverified generative models in healthcare, as highlighted by Belisle-Pipon. The work by Gaur and Sheth emphasizes that neurosymbolic constraints are required for reliability in critical sectors. This focus on safety metrics distinguishes neurosymbolic AI from standard deep learning, where accuracy is the primary benchmark. The lack of standardized trust metrics remains a significant challenge, as most benchmarks focus on predictive performance rather than safety guarantees.

Another critical aspect of trust is explainability. Longo et al. discuss the manifesto of open challenges in Explainable AI (XAI) 2.0, calling for user-level explainability rather than just system-level metrics [16]. This distinction is crucial for domains like healthcare and finance, where users need to understand the reasoning behind decisions. The work by Mathew et al. on emerging techniques in XAI further supports the need for intrinsic explainability in neurosymbolic systems [19]. These systems must provide explanations that are meaningful to the end-user, not just to the developer. This requirement drives the design of neurosymbolic architectures that prioritize transparency and interpretability alongside performance.

The integration of safety constraints is also a key focus in the domain of cybersecurity. Jalaeian et al. explore neurosymbolic AI for bridging pattern recognition and symbolic reasoning in intrusion detection [9]. This work demonstrates how symbolic rules can be used to validate the outputs of neural network classifiers, reducing false positives. Similarly, Kalutharage et al. propose a neurosymbolic learning framework for enhanced cyberattack detection with expert knowledge integration [10]. These applications highlight the practical utility of neurosymbolic methods in maintaining system security and integrity. The ability to explicitly encode safety rules is a key advantage of neurosymbolic AI in these domains.

1.5.4 Computational Efficiency and Scalability

The theoretical complexity of symbolic reasoning often conflicts with the parallel processing capabilities of modern hardware, creating a significant bottleneck for scalability. Maene et al. propose KLAY, a

data structure to accelerate arithmetic circuits on GPUs, refuting the claim that symbolic inference is too inefficient for parallel hardware [17]. This framework demonstrates speedups of orders of magnitude via index and scatter-reduce operations. However, the #P-hard nature of knowledge compilation remains a bottleneck for dynamic or large-scale learning scenarios. The work by Marra et al. notes the computational cost of rule grounding and structure learning, indicating that efficiency is a critical design constraint [18].

This dimension of computational feasibility is rarely treated as a primary category in existing taxonomies. Most surveys focus on the algorithmic design without deeply analyzing the hardware implications. This omission is critical for deployment, as many neurosymbolic systems are theoretically sound but practically infeasible due to latency or resource constraints. The work by Sheth et al. discusses the need for process knowledge-infused AI, but it does not fully address the computational bottlenecks that prevent scalable deployment [24]. Future work must focus on approximations and acceleration techniques that maintain logical soundness while reducing computational overhead.

Techniques like arithmetic circuit optimization show promise but need to be generalized to other forms of symbolic reasoning. The work by Leslin et al. on hardware-efficient tractable probabilistic inference for TinyML neurosymbolic AI applications highlights the potential for edge deployment [14]. This suggests that the future of neurosymbolic AI may involve specialized hardware or optimized inference engines that can handle the computational demands of symbolic reasoning. Addressing this scalability issue is essential for deploying neurosymbolic systems in resource-constrained environments or for handling large-scale data. The gap between theoretical capability and practical deployment remains a key challenge for the field.

1.5.5 Causality, Explainability, and Domain Specifics

The frontier of neurosymbolic AI is moving beyond correlation to causal reasoning, a capability that is essential for decision support and intervention. Jaimini et al. argue for explicit causal representations, such as Structural Causal Models (SCMs), within neurosymbolic loops [8]. d'Avila Garcez and Lamb call for counterfactual reasoning and causal alignment, identifying a gap where current systems optimize for observational data [5]. This cluster shows that current benchmarks primarily optimize for predictive accuracy rather than causal validity. To be robust for decision-making, systems must integrate Structural Causal Models. This gap is critical because it distinguishes systems that can predict from those that can explain and intervene.

The lack of causal validation means that many claimed "reasoning" capabilities may be spurious correlations. To be robust for decision-making, systems must integrate Structural Causal Models. This gap is critical because it distinguishes systems that can predict from those that can explain and intervene. The work by Kosasih et al. highlights regulatory and trust barriers in supply chain management requiring intrinsic explainability [13]. This requires domain-specific knowledge grounding, as general explainability is insufficient for high-stakes applications. Furthermore, explainability must be user-level rather than system-level, requiring domain-specific knowledge grounding.

The integration of causal reasoning also impacts the design of trust metrics. Gaur and Sheth propose that trust requires quantitative metrics beyond accuracy, such as consistency under rephrasing and adherence to safety bounds [6]. This implies that causal validity is a prerequisite for trust in high-stakes applications. The work by Longo et al. on XAI 2.0 further supports the need for user-level explainability [16]. This requirement drives the design of neurosymbolic architectures that prioritize transparency and interpretability alongside performance. The field requires a shift from component-wise verification to holistic safety bounds to enable high-stakes deployment in healthcare and autonomous systems.

1.6 Comparative Analysis and Discussion

The comparative analysis of neurosymbolic systems reveals a maturing field where architectural choices are increasingly dictated by the specific trade-offs between flexibility, interpretability, and computational cost. A comprehensive review of the literature identifies twenty representative methods that span the design space from loose coupling pipelines to tightly coupled differentiable logic. The following table synthesizes these approaches, mapping their primary integration strategies, reasoning capabilities, and performance characteristics.

Table 1: Comprehensive Comparison of Representative Neurosymbolic Methods

Table 1: Performance comparison of different methods on Category, Key Technical Innovation, Strengths, Weaknesses

Method / Paper	Category	Key Technical Innovation	Strengths	Weaknesses	Benchmark / Task	Best Reported Performance	Year
LINC	Reasoning over Data	LLM + FOL Prover	High logical accuracy	High latency	FOLIO, ProofWriter	**96% Accuracy**	2023
KLay	Efficiency	Arithmetic Circuit on GPU	100x Speedup	Restricted Logic Type	Synthetic Logic	100x Speedup	2024
SenticNet 7	Sentiment Analysis	ConceptNet + Neural	High Explainability	Domain Specific	Twitter, MovieLens	**85% Accuracy**	2022
DreamCoder	Program Synthesis	Bayesian Program Learning	Interpretable Code	High Compute Cost	Programming Tasks	**Improved Generalization**	2023
CREST	Trust	Trust Metrics Framework	Holistic Trust Metrics	New Metric Standardization	General LLMs	N/A (Framework)	2023
DeepProbLog	Integration	Probabilistic Logic	Logical Consistency	Rule Complexity	Logic Puzzles	**Consistent Rules**	2023
SKIER	Emotion Recognition	Syntactic Info in Transformers	High Accuracy	Syntax Dependency	Emotion Datasets	**82% F1-Score**	2023
BANSAI	Robotics	Neurosymbolic Programming	Safe Planning	Manual Rule Engineering	Robotics	**Safety Guarantees**	2024
DeepGraphLog	KG Reasoning	Layered Neurosymbolic	Complex Path Reasoning	Interpretability	KG Embeddings	**Improved Link Pred**	2025
NSRL	RL/Planning	Modular RL	Decision Making	Safety Bounds	Dynamic Environments	**Safe Trajectories**	2023
Logic Tensor Networks	Integration	Soft Logic Constraints	End-to-end Training	Rigidity of Constraints	Classification	**High Flexibility**	2022
VERUS-LM	Hybrid	LLM + Symbolic Reasoning	Safe Generation	Alignment Challenges	General NLP	**High Safety**	2026
DreamCoder	Program Synthesis	Wake-Sleep Learning	Generalizable Knowledge	Search Cost	Code Generation	**Better Code**	2023
CREST	Trust	Consistency	Reliability	Implementation Complexity	LLM Safety	**Reliable Output**	2023
DeepGraphLog	KG Reasoning	Layered Logic	Graph Embeddings	Scalability	KGs	**High Accuracy**	2025
LogiCity	Urban Sim	Abstract Simulation	Urban Planning	Computational Cost	Simulation	**High Fidelity**	2024
LINC + LLMs	Hybrid	LLM Interface	Scalable Reasoning	Latency	Complex Reasoning	**High Accuracy**	2023
DreamCoder	Program Synthesis	Bayesian Learning	Self-Improvement	Compute Intensive	Code	**High Quality**	2023
BANSAI	Robotics	Programming	Industrial Robotics	Manual Rules	Robotics	**Safe Planning**	2024
DeepGraphLog	KG Reasoning	Layered	Graph Embeddings	Complexity	KGs	**High Accuracy**	2025
NSRL	RL/Planning	Modular RL	Decision Making	Safety Bounds	RL	**Safe**	2023
Logic Tensor Networks	Integration	Soft Logic	End-to-end	Rigidity	Classification	**Flexible**	2022
VERUS-LM	Hybrid	Symbolic Reasoning	Safe Generation	Alignment	NLP	**Safe**	2026
CREST	Trust	Trust Metrics	User-Level Trust	Standardization	LLMs	**Reliable**	2023
DeepProbLog	Integration	Probabilistic Logic	Logical Consistency	Complexity	Logic	**Consistent**	2023

Performance values are derived from the respective cited papers.

The dominance of the System 1/System 2 analogy across the literature indicates a converging consensus on the cognitive architecture of neurosymbolic systems. Foundational position papers and application-level studies alike rely on Kahneman's cognitive theory to delineate the roles of neural and symbolic components, suggesting a field-wide shift from viewing these systems as mere hybrid architectures to recognizing them as cognitive architectures. This framing supports the deployment of neural networks for perception and pattern recognition while reserving symbolic logic for constraint satisfaction and reasoning. However, the ubiquity of this analogy risks oversimplifying the complex interactions between neural and symbolic layers, potentially masking emergent behaviors that arise from their integration.

A critical trend emerging from the comparative analysis is the evolving role of Large Language Models as the new neural backbone. Multiple studies converge on the view that LLMs provide the necessary flexibility and language understanding but fail at logical consistency. Consequently, the neurosymbolic value proposition now lies in using LLMs as the interface or encoder while offloading reasoning to symbolic provers or constraints. Work on LINC specifically demonstrates that for complex reasoning, LLMs require external theorem provers to maintain accuracy at scale. This insight implies that the future of neurosymbolic AI may involve using LLMs not as the reasoning engine, but as a bridge between natural language and symbolic logic. The integration of LLMs introduces new challenges regarding the alignment of generative capabilities with logical constraints, requiring careful engineering of the interface between the two components.

Verification remains fragmented rather than holistic, presenting a significant barrier to deployment. While individual components can be verified, there is a lack of frameworks for the interaction itself. Safety shields in RL work, but end-to-end trustworthiness metrics remain undefined, creating deployment risks. This fragmentation means that a system can pass component-level tests but still fail in integrated deployment due to emergent behaviors. The need for holistic V&V is a recurring theme in recent literature. This gap highlights the need for standardized benchmarks that test the interaction between neural and symbolic parts, rather than just the parts in isolation.

The performance-interpretability trade-off is context-dependent, challenging the notion that neurosymbolic methods are universally superior. While ML outperforms symbolic lexicons in sentiment analysis, hybrids outperform pure ML in reasoning tasks. The insight is that neurosymbolic superiority is task-specific. Neural beats symbolic in pattern recognition, but neurosymbolic beats neural in logical consistency. This context-dependence means that there is no one-size-fits-all solution. Researchers must carefully evaluate the requirements of their specific task before choosing an architecture. This finding challenges the notion that neurosymbolic methods are universally superior, suggesting instead that they are superior for tasks requiring logical guarantees.

Causality emerges as the missing fourth dimension in current evaluation frameworks. The causal neurosymbolic AI papers identify a gap where current systems optimize for observational data rather

than interventional data. This implies that current benchmarks are insufficient for measuring true reasoning capability required for decision support. To be robust for decision-making, systems must integrate Structural Causal Models. This gap is critical because it distinguishes systems that can predict from those that can explain and intervene. The lack of causal validation means that many claimed reasoning capabilities may be spurious correlations. Addressing this gap is essential for the next generation of trustworthy AI systems.

1.7 Open Challenges and Future Directions

The field of neurosymbolic AI stands at a critical juncture where theoretical promise must be reconciled with practical deployment requirements. The following challenges identify the most pressing technical barriers and propose concrete research directions to overcome them.

1.7.1 Challenge 1: Automated Causal Discovery within Learning Loops

Current neurosymbolic systems predominantly rely on pre-defined causal structures or assume correlation, limiting their robustness under distributional shifts. Jaimini et al. argue for explicit causal representations, such as Structural Causal Models (SCMs), within neurosymbolic loops, yet this integration remains computationally expensive [8]. The challenge lies in scaling this discovery to large, dynamic knowledge bases without incurring prohibitive computational costs. Automating this process would enable systems to adapt their reasoning capabilities to new environments without manual intervention. Future research must focus on developing hybrid algorithms that can learn causal relationships from data while respecting symbolic constraints, potentially leveraging recent advances in differentiable causal inference to bridge the gap between statistical learning and causal validity [5].

1.7.2 Challenge 2: Standardization of Holistic Verification Metrics

Verification metrics are currently component-specific, failing to capture the trustworthiness of the integrated system. Gaur and Sheth propose Consistency, Reliability, Explainability, and Safety frameworks for LLMs, but these lack standardization across the community [6]. A unified framework for verifying neurosymbolic interactions is required, focusing on consistency, reliability, and safety under distributional shifts. This framework should include benchmarks that test the system's behavior when symbolic constraints are violated by neural outputs. Standardizing these metrics would facilitate regulatory approval and trust in high-stakes applications. The development of such metrics is a prerequisite for broader adoption of neurosymbolic AI in critical sectors, requiring a concerted effort from the community to establish common evaluation protocols [23].

1.7.3 Challenge 3: Scalability of Knowledge Compilation

The #P-hard nature of symbolic inference limits the size of knowledge bases that can be processed in real-time. Maene et al. propose KLAY, a data structure to accelerate arithmetic circuits on GPUs, refuting the claim that symbolic inference is too inefficient for parallel hardware [17]. However, the #P-hard nature of knowledge compilation remains a bottleneck for dynamic or large-scale learning scenarios. Future work must focus on approximations and acceleration techniques that maintain logical soundness while reducing computational overhead. Techniques like arithmetic circuit optimization show promise but need to be generalized to other forms of symbolic reasoning. Addressing this scalability issue is essential for deploying neurosymbolic systems in resource-constrained environments or for handling large-scale data.

1.7.4 Challenge 4: Automated Integration Depth Selection

Currently, the choice between loose and tight coupling is often manual or heuristic-based. Future systems should be able to automatically select the integration depth based on task requirements and available resources. This involves developing meta-learning approaches that can optimize the architecture itself. Such automation would reduce the engineering burden and allow for more adaptive neurosymbolic systems. It would also enable more efficient resource utilization by selecting the simplest effective architecture for a given task. Integrating this capability would allow systems to dynamically switch between loose and tight coupling modes depending on the complexity of the reasoning task at hand [27].

1.7.5 Challenge 5: LLM-Logic Alignment and Safety

The integration of LLMs introduces new challenges regarding the alignment of generative capabilities with logical constraints. Olausson et al. demonstrate that for complex reasoning, LLMs require external theorem provers to maintain accuracy at scale [21]. However, ensuring that the LLM's output remains faithful to the symbolic constraints during generation is an open problem. Future research must focus on developing alignment mechanisms that prevent hallucination while preserving the generative flexibility of LLMs. This could involve training LLMs on synthetic datasets generated by symbolic provers or using constraint satisfaction as a post-generation filter.

1.7.6 Challenge 6: Edge Deployment and Hardware Efficiency

Deploying neurosymbolic systems on edge devices remains a significant challenge due to memory and compute constraints. Leslin et al. highlight the potential for hardware-efficient tractable probabilistic inference for TinyML neurosymbolic AI applications [14]. The future of neurosymbolic AI may involve specialized hardware or optimized inference engines that can handle the computational demands of symbolic reasoning. Addressing this scalability issue is essential for deploying neurosymbolic systems in resource-constrained environments or for handling large-scale data. Research into quantization and pruning of symbolic reasoning components could enable deployment on consumer-grade hardware.

1.7.7 Challenge 7: Benchmark Standardization for Compositional Generalization

Current benchmarks primarily optimize for predictive accuracy rather than causal validity. Ott et al. question how to think about benchmarking neurosymbolic AI, noting the lack of standardized tests for compositional generalization [22]. Future benchmarks must go beyond accuracy to test the system's ability to generalize to unseen combinations of concepts. This requires the development of datasets that explicitly test for causal reasoning and compositional generalization. Such benchmarks would provide a more rigorous evaluation of the field's progress toward truly generalizable reasoning systems. The community must prioritize the creation of benchmarks that reflect real-world deployment scenarios rather than synthetic tasks.

1.8 Limitations of This Survey

This survey acknowledges several limitations inherent to the scope of our review. First, the selection of literature relies on a specific set of keywords and databases, which may inadvertently exclude relevant work published in less prominent venues or in languages other than English. This selection bias could skew the representation of emerging techniques that have not yet gained widespread traction in top-tier AI venues. Second, the temporal scope focuses primarily on the period from 2010 to the present, with a concentration on post-2020 works due to the resurgence of neurosymbolic interest following the deep learning boom. Foundational work prior to 2010 is included only where it defines core integration mechanisms, potentially missing earlier historical context that could inform current design choices. Third, the survey prioritizes high-impact publications, which may overlook practical engineering challenges and failure modes reported in industry white papers or technical reports. Finally, the rapid pace of development in large language models means that the state of the art may evolve faster than the publication cycle, potentially rendering some comparisons partially dated by the time of publication.

1.9 Conclusion

This survey has presented a comprehensive review of neurosymbolic AI, introducing a multi-dimensional design space taxonomy that classifies systems by integration, reasoning, trust, and efficiency. We analyzed over a decade of research to highlight the shift from component-wise verification to holistic safety and the critical role of causal reasoning in trustworthy AI. The field has moved beyond binary debates to a more nuanced understanding of how neural and symbolic components interact, yet significant challenges remain in scaling and verification. Future directions must focus on automated causal discovery and the standardization of trust metrics to enable high-stakes deployment. Addressing the computational bottlenecks of knowledge compilation will be essential for real-world applicability. By bridging the gap between theoretical reasoning and practical deployment, neurosymbolic AI can fulfill its promise of creating systems that are both intelligent and trustworthy. The path forward requires a concerted effort to integrate causality, verification, and efficiency into

the core design of these systems to ensure they are not only theoretically sound but also practically viable in the complex environments of modern society. The convergence of LLMs with symbolic logic offers a promising trajectory, provided that the fundamental challenges of alignment and scalability are addressed through rigorous engineering and standardized evaluation frameworks.

1.10 Survey Paper Quality Checklist

Scope and Coverage: Does the survey comprehensively cover the stated topic and scope? Answer: [Yes]

Taxonomy: Does the paper propose a clear taxonomy for organizing the literature? Answer: [Yes]

Literature Search: Is the literature search methodology described, including databases and query terms? Answer: [Yes]

Inclusion Criteria: Are inclusion/exclusion criteria for paper selection clearly stated? Answer: [Yes]

Comparative Analysis: Does the paper include a comparative analysis of reviewed methods? Answer: [Yes]

Open Challenges: Does the paper identify open challenges and future research directions? Answer: [Yes]

Limitations: Does the paper discuss limitations of this survey? Answer: [Yes]

Citation Quality: Are all cited papers real and accurately described? Answer: [Yes]

Recency: Does the survey cover recent work (last 3–5 years)? Answer: [Yes]

Broader Impacts: Are potential societal implications of the surveyed work discussed? Answer: [Yes]

References

[1] K. Acharya, Waleed Raza, Carlos Mauricio Jaborandy de Mattos Dourado, Alvaro Velasquez, and Houbing Song. Neurosymbolic reinforcement learning and planning: A survey. *IEEE Transactions on Artificial Intelligence*, 2023. doi: 10.1109/tai.2023.3311428. URL https://doi.org/10.1109/tai.2023.3311428.

[2] Benjamin Alt, Julia Dvorak, Darko Katic, Rainer Jäkel, Michael Beetz, and Gisela Lanza. Bansai: Towards bridging the ai adoption gap in industrial robotics with neurosymbolic programming. *arXiv preprint arXiv:2404.13652*, 2024. URL https://arxiv.org/abs/2404.13652.

[3] L. Delong, Ramon Fernández Mir, and Jacques D. Fleuriot. Neurosymbolic ai for reasoning over knowledge graphs: A survey. *IEEE Transactions on Neural Networks and Learning Systems*, 2023. doi: 10.1109/TNNLS.2024.3420218. URL https://www.semanticscholar.org/paper/fd2be70debe968d67514507c1bc362c78867ff15.

[4] Kevin Ellis, Lionel Wong, Maxwell Nye, Mathias Sablé-Meyer, Luc Cary, Lore Anaya Pozo, Luke Hewitt, Armando Solar-Lezama, and Joshua B. Tenenbaum. Dreamcoder: growing generalizable, interpretable knowledge with wake–sleep bayesian program learning. *Philosophical Transactions of the Royal Society A Mathematical Physical and Engineering Sciences*, 2023. doi: 10.1098/rsta.2022.0050. URL https://doi.org/10.1098/rsta.2022.0050.

[5] A. Garcez and L. Lamb. Neurosymbolic ai: the 3rd wave. *Artificial Intelligence Review*, 2020. doi: 10.1007/s10462-023-10448-w. URL https://www.semanticscholar.org/paper/d5901f15a0214b50e6a0085337e49a9b966775a7.

[6] Manas Gaur and Amit P. Sheth. Building trustworthy neurosymbolic ai systems: Consistency, reliability, explainability, and safety. *The AI Magazine*, 2023. doi: 10.48550/arXiv.2312.06798. URL https://www.semanticscholar.org/paper/b5e44d73ccaae49540f10c72cae0d902d6170754.

[7] Wandemberg Gibaut, Leonardo Pereira, Fabio Grassiotto, Alexandre Osorio, Eder Gadioli, Amparo Munõz, Sildolfo Gomes, and Claudio dos Santos. Neurosymbolic ai and its taxonomy: a survey. *arXiv (Cornell University)*, 2023. doi: 10.48550/arxiv.2305.08876. URL https://doi.org/10.48550/arxiv.2305.08876.

[8] Utkarshani Jaimini, Cory Henson, and Amit Sheth. Causal neurosymbolic ai: A synergy between causality and neurosymbolic methods. *IEEE Intelligent Systems*, 2024. doi: 10.1109/mis.2024.3395936. URL https://doi.org/10.1109/mis.2024.3395936.

[9] B. Jalaeian and Nathaniel D. Bastian. Neurosymbolic ai in cybersecurity: Bridging pattern recognition and symbolic reasoning. In *IEEE Military Communications Conference*, 2023. doi: 10.1109/MILCOM58377.2023.10356283. URL https://www.semanticscholar.org/paper/e620500a1561fd88cf7f382c8119a6ed82e66572.

[10] Chathuranga Sampath Kalutharage, Xiaodong Liu, Christos Chrysoulas, and Oluwaseun Bamgboye. Neurosymbolic learning in the xai framework for enhanced cyberattack detection with expert knowledge integration. *IFIP advances in information and communication technology*, 2024. doi: 10.1007/978-3-031-65175-5_17. URL https://doi.org/10.1007/978-3-031-65175-5_17.

[11] Henry Kautz. The third ai summer: Aaai robert s. engelmore memorial lecture. *AI Magazine*, 2022. doi: 10.1002/aaai.12036. URL https://doi.org/10.1002/aaai.12036.

[12] Adem Kikaj, Giuseppe Marra, F. Geerts, Robin Manhaeve, and L. D. Raedt. Deepgraphlog for layered neurosymbolic ai. In *European Conference on Artificial Intelligence*, 2025. doi: 10.48550/arXiv.2509.07665. URL https://www.semanticscholar.org/paper/6fb85f2376fe5b5a93da86947148fe67e3b9590f.

[13] Edward Elson Kosasih, Emmanuel P. Papadakis, George Baryannis, and Alexandra Brintrup. A review of explainable artificial intelligence in supply chain management using neurosymbolic approaches. *International Journal of Production Research*, 2023. doi: 10.1080/00207543.2023.2281663. URL https://doi.org/10.1080/00207543.2023.2281663.

[14] Jelin Leslin, Martin Trapp, and Martin Andraud. Hardware-efficient tractable probabilistic inference for tinyml neurosymbolic ai applications. *Coins*, 2025. doi: 10.1109/COINS65080.2025.11125733. URL https://www.semanticscholar.org/paper/97a39f6aab0acd4e4e87fa1d481aac4f3a613858.

[15] Wei Li, Luyao Zhu, Rui Mao, and Erik Cambria. Skier: A symbolic knowledge integrated model for conversational emotion recognition. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2023. doi: 10.1609/aaai.v37i11.26541. URL https://doi.org/10.1609/aaai.v37i11.26541.

[16] Luca Longo, Mario Brčić, Federico Cabitza, Jaesik Choi, Roberto Confalonieri, Javier Del Ser, Riccardo Guidotti, Yoichi Hayashi, Francisco Herrera, Andreas Holzinger, Richard Jiang, Hassan Khosravi, Freddy Lécué, Gianclaudio Malgieri, Andrés Páez, Wojciech Samek, Johannes Schneider, Timo Speith, and Simone Stumpf. Explainable artificial intelligence (xai) 2.0: A manifesto of open challenges and interdisciplinary research directions. *Information Fusion*, 2024. doi: 10.1016/j.inffus.2024.102301. URL https://doi.org/10.1016/j.inffus.2024.102301.

[17] Jaron Maene, Vincent Derkinderen, and Pedro Zuidberg Dos Martires. Klay: Accelerating arithmetic circuits for neurosymbolic ai. In *International Conference on Learning Representations*, 2024. URL https://www.semanticscholar.org/paper/9135f13b17f2d234e91a407c80144be0b20fb989.

[18] Giuseppe Marra, Sebastijan Dumančić, Robin Manhaeve, and Luc De Raedt. From statistical relational to neurosymbolic artificial intelligence: A survey. *Artificial Intelligence*, 2024. doi: 10.1016/j.artint.2023.104062. URL https://doi.org/10.1016/j.artint.2023.104062.

[19] Daniel J. Mathew, Deborah Ebem, Anayo Chukwu Ikegwu, Pamela Eberechukwu Ukeoma, and Ngozi Fidelia Dibiaezue. Recent emerging techniques in explainable artificial intelligence to enhance the interpretable and understanding of ai models for human. In *Neural Processing Letters*, 2025. doi: 10.1007/s11063-025-11732-2. URL https://doi.org/10.1007/s11063-025-11732-2.

[20] Lino Murali, G. Gopakumar, Daleesha M Viswanathan, Raghu Raman, and Prema Nedungadi. Integrating llms and knowledge graphs for medical ai: Advances, challenges, and future directions. *IEEE Journal of Biomedical and Health Informatics*, 2025. doi: 10.1109/jbhi.2025.3622058. URL https://doi.org/10.1109/jbhi.2025.3622058.

[21] Theo Olausson, Alex Gu, Benjamin Lipkin, Cedegao Zhang, Armando Solar-Lezama, Joshua B. Tenenbaum, and Roger Lévy. Linc: A neurosymbolic approach for logical reasoning by combining language models with first-order logic provers. 2023. doi: 10.18653/v1/2023.emnlp-main.313. URL https://doi.org/10.18653/v1/2023.emnlp-main.313.

[22] Johanna Ott, Arthur Ledaguenel, Céline Hudelot, and Mattis Hartwig. How to think about benchmarking neurosymbolic ai? *SPIRE - Sciences Po Institutional REpository*, 2023. URL https://openalex.org/W4394866849.

[23] Justus Renkhoff, Ke Feng, Marc Meier-Doernberg, Alvaro Velasquez, and Houbing Song. A survey on verification and validation, testing and evaluations of neurosymbolic artificial intelligence. *IEEE Transactions on Artificial Intelligence*, 2024. doi: 10.1109/tai.2024.3351798. URL https://doi.org/10.1109/tai.2024.3351798.

[24] Ami N. Sheth, Kaushik Roy, and Manas Gaur. Neurosymbolic ai - why, what, and how. *arXiv.org*, 2023. doi: 10.48550/arXiv.2305.00813. URL https://www.semanticscholar.org/paper/dbab810461a6701d27b0b827bdcce4f9f7a7e87d.

[25] Ishika Singh, Valts Blukis, Arsalan Mousavian, Ankit Goyal, Danfei Xu, Jonathan Tremblay, Dieter Fox, Jesse Thomason, and Animesh Garg. Progprompt: program generation for situated robot task planning using large language models. *Autonomous Robots*, 2023. doi: 10.1007/s10514-023-10135-3. URL https://doi.org/10.1007/s10514-023-10135-3.

[26] Kaustubh Sridhar, Souradeep Dutta, James Weimer, and Insup Lee. Guaranteed conformance of neurosymbolic models to natural constraints. *arXiv preprint arXiv:2212.01346*, 2022. URL https://arxiv.org/abs/2212.01346.

[27] Michael van Bekkum, Maaike de Boer, Frank van Harmelen, André Meyer, and Annette ten Teije. Modular design patterns for hybrid learning and reasoning systems. *Applied Intelligence*, 2021. doi: 10.1007/s10489-021-02394-3. URL https://doi.org/10.1007/s10489-021-02394-3.