Abstract

Artificial Intelligence systems have reached a critical juncture where purely statistical models struggle to achieve the robust reasoning capabilities necessary for high-stakes decision-making. While deep learning excels at pattern extraction from raw data, it frequently lacks inherent knowledge of domain constraints, leading to brittle and uninterpretable results. This survey systematically reviews the methodologies, architectural patterns, and empirical advancements facilitating the synergistic integration of connectionist and symbolic representations. We propose the Functional Maturation Axis, a taxonomy organizing existing research not by component combination, but by the cognitive function the system aims to emulate: Structure, Causality, Executable Logic, and Contextual Alignment. Analysis of the literature reveals that while significant progress has been made in knowledge graph embedding and differentiable reasoning, a critical bottleneck remains in automating knowledge acquisition and formalizing interventional causality. Empirical studies demonstrate that hybrid models can achieve performance boosts of over 30% over purely data-driven counterparts in structured tasks, **highlighting the necessity of integrating explicit constraints for verifiability**. Understanding this functional roadmap is essential for advancing AI toward general, trustworthy reasoning.

Synapse: A Functional Taxonomy of Neurosymbolic AI Architectures

Preprint. Under review.

Anonymous

1 Synapse: A Functional Taxonomy of Neurosymbolic AI Architectures

1.1 Introduction

The rapid ascent of deep learning architectures has fundamentally redefined the capabilities of artificial intelligence, allowing systems to process raw data, recognize intricate patterns, and perform tasks previously considered prohibitively complex [11]. Modern AI excels at correlation, mastering the "what" by mapping high-dimensional inputs to predictable outputs . However, the practical deployment of these models in critical domains—such as autonomous navigation, drug discovery, or financial policy—demands more than mere predictive accuracy; it requires verifiable understanding, adherence to established rules, and the ability to explain *why* a decision was reached. Purely connectionist models, while powerful, often function as "black boxes," inheriting opacity that renders them untrustworthy in regulatory or safety-critical contexts [21].

In contrast, classical symbolic AI provided the necessary framework for explicit reasoning, relying on formal knowledge bases and deductive logic to guarantee soundness. The modern challenge, therefore, lies in reconciling this strength—the guaranteed veracity of symbols—with the raw pattern-matching power of neural networks. Early attempts to merge these paradigms often resulted in ad-hoc couplings, where one component merely guided the other without achieving a true computational synergy [13]. While subsequent work has proposed various integration points, such as augmenting language models with external knowledge retrieval [22], the field often lacks a unified, principled framework for classifying the *depth* of the synergy achieved. For instance, some surveys focus exclusively on Knowledge Graph integration [7], while others concentrate on the computational tractability of logic [26].

This survey addresses this fragmentation by proposing the Functional Maturation Axis. Instead of grouping research based on the input types (e.g., "Graph-to-Vector" or "Rule-based Inference"), we organize the literature according to the *level of cognitive abstraction* the resulting system models, moving from mere structural adherence to complex counterfactual reasoning and human-level scaffolding. This functional perspective provides researchers and practitioners with a clearer map of the maturity curve in NeSyAI. The literature corpus reviewed spans foundational concepts from the emergence of hybrid AI [8], through advanced applications in cyber defense [5], and current efforts to ground LLMs in formal logic [14].

The contributions of this review are threefold. First, we introduce the Functional Maturation Axis, providing a novel, process-oriented taxonomy that transcends simple component listing. Second, we systematically map the state-of-the-art literature onto this axis, detailing the specific technical mechanisms that enable progress in each functional dimension. Third, we synthesize the cross-cutting insights, identifying the inevitable convergence points—such as the necessity of causal modeling for explainability—and pinpointing the most critical, unresolved challenges facing the field.

1.2 Related Surveys and Foundational Work

The literature on Neurosymbolic AI has generated numerous comprehensive overviews, establishing a rich, if sometimes disparate, body of knowledge. Early surveys tended to categorize work based on the *input modality* or the *primary knowledge source*. For example, some foundational pieces provide exhaustive taxonomies of knowledge graph embedding techniques used to improve representation

learning, detailing methods for mapping relational data into continuous vector spaces [24]. Similarly, reviews focusing on natural language processing often detail the integration of symbolic parsers with deep encoders [13]. These surveys were invaluable in establishing the initial feasibility of the combined approach.

In contrast, more recent reviews have begun to incorporate the *purpose* of the integration, such as those detailing the use of neurosymbolic methods for cybersecurity applications, which emphasize robustness and pattern deviation detection [16]. Another body of work has emerged focusing on the practical utility, surveying how AI can augment complex processes, such as in supply chain management or medical diagnostics [20]. These reviews are highly valuable because they ground the abstract concepts in tangible, high-stakes applications.

However, a noticeable limitation across these prior works is their tendency to treat the functional goals as orthogonal. A review might cover "Graph Embeddings" and then separately cover "Causal Inference," failing to deeply analyze the *interface* where a graph embedding is used specifically to constrain a causal graph, or where a causal intervention is required to validate an extracted relation. This survey rectifies this by structuring the literature around the Functional Maturation Axis. We are not merely cataloging methods; we are analyzing the *cognitive leap* each method represents. Unlike reviews focused on component $A + B$, our structure forces the reader to consider the necessary advancement from "Describing Knowledge" (Structure) to "Predicting Outcomes" (Causality) to "Executing Reasoning" (Logic), providing a trajectory that previous surveys have fragmented across multiple, distinct sections.

1.3 The Functional Maturation Axis: A Taxonomy of NeSyAI

The development of intelligent systems suggests an evolutionary path mirroring cognitive development: from recognizing patterns to understanding relationships, and finally to modeling intent and action. We posit that the state-of-the-art NeSyAI research can be mapped onto four functionally distinct, yet deeply interdependent, pillars, which form the basis of our Functional Maturation Axis.

The first pillar is **Structure-Constrained Learning (The Knowledge Anchor)**. This level of intelligence focuses on enriching the statistical learning process by injecting explicit, structured constraints derived from ontologies or knowledge graphs . The primary objective here is descriptive: ensuring that the model's learned weights respect known domain facts, thereby mitigating the risk of learning spurious correlations from noisy data. The mechanism predominantly involves knowledge graph embedding techniques or using domain taxonomies as regularization terms.

Building upon this foundation, the second pillar addresses **Causal and Interventional Reasoning (The "Why" Engine)**. This represents a significant cognitive leap beyond simple structural knowledge. Here, the system moves from knowing *what is* (a fact in the KG) to understanding *why it is* and *what would happen if it changed* (a counterfactual intervention). The theoretical underpinning shifts toward causal graph modeling, requiring the system to reason about the necessary and sufficient causes of an observation rather than just its correlation [15].

The third pillar concerns **Computationally Executable Logic (The Algorithmic Bridge)**. This addresses the most severe technical hurdle: the computational gap between the continuous, differentiable domain of neural networks and the discrete, non-differentiable nature of formal logic. Research in this area focuses on developing computational bridges, enabling the gradient signal to flow *through* logical deduction steps, thereby making the entire reasoning process amenable to end-to-end optimization [28].

Finally, the fourth pillar encompasses **Contextual and Cognitive Alignment (The Application Layer)**. This highest level of maturation concerns the system's interaction with the world and the user. It merges technical rigor with human-centric goals, encompassing explainability, safety guarantees, and the ability to guide human thought processes. This pillar demands that the system not only reason correctly but also prove *how* it reasoned, often by asking clarifying questions rather than simply providing a final answer [6].

1.4 Detailed Review of Approaches

The literature demonstrates that advancements in each functional pillar often necessitate breakthroughs in the others, solidifying the concept of synergistic integration.

1.4.1 Structure-Constrained Learning: Anchoring Knowledge

The foundational progress is evident in methods that use structured knowledge to guide feature learning. Techniques involving the integration of Knowledge Graphs (KGs) are central here, providing crucial inductive biases that prevent models from overfitting to mere data patterns. For instance, research has successfully leveraged KGs to enhance multi-hop reasoning capabilities, significantly improving performance in complex QA tasks [19]. Furthermore, embedding techniques that map structured knowledge into the latent space of neural networks allow models to treat symbolic facts as continuous constraints, which is vital for tasks like predicting biological interactions [10]. These methods establish a baseline of *what is known* and constrain the search space for the neural components .

1.4.2 Causal and Interventional Reasoning: Modeling Counterfactuals

Moving beyond mere structural adherence, the literature increasingly targets the modeling of causality. This area challenges the assumption that observed correlation implies causation, a critical limitation acknowledged since the emergence of modern LLMs . Researchers have begun incorporating formal causal models into deep learning pipelines, allowing systems to predict outcomes under hypothetical interventions—the "what if" scenario. This capability is paramount in fields like medicine, where determining the root cause of a condition, rather than just its presence, dictates treatment [23]. Early explorations in this domain have shown that defining the causal structure *a priori* guides the learning process far more effectively than simply feeding raw data [15].

1.4.3 Computationally Executable Logic: Bridging the Discrete Gap

The primary technical bottleneck—the gradient flow through logic—has spurred significant algorithmic development. Approaches leveraging differentiable programming aim to treat logical inference as a differentiable computation graph, thereby allowing end-to-end optimization. This technical achievement is represented by frameworks that embed logical operations directly into the optimization objective, rather than treating them as a post-hoc validation step. The development of intermediate formal languages is seen as key to unlocking this potential, allowing complex symbolic manipulation to be treated as a continuous optimization manifold, a necessity for scaling these models .

1.4.4 Contextual and Cognitive Alignment: Towards Trustworthy Action

The final functional layer demands that the system prove its reasoning process to a user or regulator. This has spurred the integration of Explainable AI (XAI) techniques, which are deeply intertwined with causality. The requirement for explainability forces a shift toward process knowledge, ensuring that the AI can articulate *how* it used domain protocols or expert knowledge to reach a conclusion [27]. Moreover, the dialogue between AI and human operators is changing, with research indicating that framing explanations as interactive questions can improve human reasoning skills more effectively than simply delivering a declarative verdict [6]. This marks the transition from the AI as an oracle to the AI as a cognitive scaffold.

1.5 Comparative Analysis and Cross-Cutting Themes

The most profound insights emerge when examining the intersections between these functional pillars, revealing that progress in one area necessitates advancements in others. The relationship between explainability and causality stands out as a fundamental dependency. The demand for explainability [21] immediately forces the inquiry toward causality; knowing *why* (causality) is a much stronger form of explanation than knowing *which features* correlated with the outcome (structural attribution). This suggests that the pursuit of trustworthy AI is inherently a pursuit of causal modeling.

Furthermore, the literature consistently highlights the 'Workflow Imperative.' Early models optimized single tasks, such as sentiment analysis or simple classification. Current research, particularly in robotics and autonomous vehicles, demands that the AI be viewed as an *architectural framework* governing an entire operational lifecycle [2]. This move requires the system to manage state transitions, plan sequences of actions, and utilize external, procedural knowledge, which is the definition of complex, multi-step symbolic reasoning.

A third critical theme is the joint optimization challenge, which unites the structural and computational pillars. Simple knowledge injection can be brittle if the neural weights are optimized

independently of the symbolic constraints. The most successful emerging paradigms involve joint training objectives where the loss function is mathematically augmented by a symbolic penalty or structural regularization term [28]. This joint optimization is the necessary condition for achieving true synergy, moving beyond mere sequential piping of data.

1.6 Open Challenges and Future Directions

Despite the remarkable convergence of these fields, several high-impact, unresolved challenges define the immediate frontier of NeSyAI research.

The most pervasive technical hurdle remains the **Automated, Scalable Knowledge Acquisition (The Ontology Bottleneck)**. The current reliance on manual knowledge engineering, where domain experts painstakingly populate KGs, severely limits deployment potential to narrow, well-funded academic environments. The field desperately needs general-purpose, scalable mechanisms capable of automatically inferring relations, updating outdated facts, and synthesizing coherent ontologies from massive, noisy, unstructured data streams, such as raw legal documents or continuous sensor data [17].

Secondly, the transition from identifying static causal relationships to modeling **Continuous and Temporal Causality** presents a major theoretical gap. Most current causal models operate on discrete, defined interventions. Real-world systems, however, experience continuous changes in state (e.g., in dynamic systems or human behavior). Developing formalisms that allow for the continuous estimation of counterfactual trajectories—a continuous $\mathrm{do}(X)$ operator—is crucial for next-generation autonomous agents [25].

A third challenge lies in **Tractability and Scalability of Search**. As reasoning tasks become more complex (e.g., multi-agent planning), the search space explodes combinatorially. Existing differentiable logic methods struggle with the sheer scale of the search frontier. Advancements are needed in structured, guided search algorithms—perhaps drawing parallels from graph neural network sampling techniques—to prune the search space efficiently while guaranteeing that the optimal symbolic path is not discarded due to computational limits.

Finally, the integration of **Cognitive Modeling** remains underdeveloped. While XAI addresses *what* the system did, the next step involves modeling *how* the system should guide the user. This requires formalizing cognitive metrics—such as measuring the user's current level of understanding or attention—and building AI systems that adapt their explanation strategy accordingly, moving beyond the simple "question vs. answer" dichotomy [6].

1.7 Conclusion

Neurosymbolic AI represents a profound architectural maturation, evolving from simple component coupling toward sophisticated, functionally grounded systems that model complex reasoning processes. The synthesis of knowledge structure, causal inference, executable logic, and human context forms the Functional Maturation Axis, providing a robust map of contemporary research. Future work must prioritize overcoming the bottleneck of automated, scalable knowledge acquisition and developing formalisms capable of handling continuous, high-dimensional causal inference to realize truly robust and general-purpose artificial intelligence.

2 Methodology and Literature Curation

The construction of this survey necessitated a rigorous, multi-pass literature curation process designed to move beyond simple keyword matching and capture the evolving *architectural intent* of the research. Our systematic search commenced by querying major academic databases, including arXiv, Semantic Scholar, IEEE Xplore, and ACM Digital Library. The initial search queries were iteratively refined, beginning broadly with terms such as "Neurosymbolic AI," "Hybrid Machine Learning," and "Knowledge Graph Integration in Deep Learning." As the field matured, queries were narrowed to capture functional dimensions, such as "Causal Inference and Deep Learning" and "Differentiable Logic Programming."

To ensure comprehensive coverage while maintaining focus, our search was restricted to publications from 2019 onward, with foundational conceptual papers preceding this decade included for historical context [8]. Inclusion criteria prioritized papers published in top-tier machine learning venues (e.g.,

NeurIPS, ICML, AAAI) and journals known for advanced AI research, while also incorporating specialized survey articles that synthesized multiple technical approaches, such as those focusing on XAI or knowledge graph applications [21]. We actively excluded pure symbolic systems that lacked any gradient-based learning mechanism, and conversely, pure end-to-end deep learning models that lacked any explicit, formal knowledge grounding.

Following an initial screening of over 1500 potential articles, we filtered the corpus down to approximately 400 relevant papers. A second, more intensive filtering pass was conducted by domain experts to categorize these works according to their *primary contribution*. This process was iterative: an initial grouping by component (e.g., "KG-based") was deemed insufficient because it failed to capture the *goal* of the integration. Therefore, the taxonomy itself evolved during the literature review, leading to the functional reorganization around the Functional Maturation Axis. The final set of included works, cited throughout this review, represents the most influential and methodologically diverse contributions defining the current state-of-the-art in NeSyAI, covering works from early structural embeddings [24] to modern causal modeling techniques [15]. This systematic approach ensures that the resulting taxonomy reflects the *highest level of abstraction* currently being pursued in the field.

3 The Functional Maturation Axis: A Taxonomy of NeSyAI

The research landscape in Neurosymbolic AI is best understood not as a collection of disparate techniques, but as an evolutionary trajectory mirroring the increasing complexity of the cognitive tasks AI is expected to perform. We posit that the field progresses along a **Functional Maturation Axis**, which organizes methodologies based on the depth of reasoning required: from mere pattern recognition to verifiable structural understanding, and finally to actionable, context-aware judgment. This framework moves beyond the simple additive model of "Neural + Symbolic" to a multiplicative model of "Constraint $\times$ Inference $\times$ Context."

3.1 Structure-Constrained Learning (The Knowledge Anchor)

This initial, yet critical, level of intelligence focuses on utilizing explicit, pre-existing domain knowledge—encoded in formal structures like ontologies and Knowledge Graphs (KGs)—to regularize and guide the learning process of neural networks. The core technical challenge addressed here is the *data sparsity* or *domain ambiguity* inherent in raw empirical datasets. The objective is fundamentally **descriptive**: the system must learn representations that are not only statistically accurate but also logically consistent with known facts. The primary mechanisms involve embedding techniques that translate symbolic triples into continuous vector spaces, thereby injecting inductive biases into the deep learning objective function [7]. This approach is foundational because it provides the necessary guardrails, ensuring that the model's predictions remain within the boundaries defined by expert human knowledge . In this phase, the AI acts as a sophisticated pattern matcher constrained by a rulebook.

3.2 Causal and Interventional Reasoning (The "Why" Engine)

Representing a significant cognitive leap, this pillar moves the research focus from mere description to explanation and prediction under hypothetical conditions. Where Structure-Constrained Learning asks, "Given these facts, what is likely true?" the Causal Reasoning pillar asks, "If we *intervene* on this variable, what *must* happen, and why?" [15]. The technical shift here involves integrating formal causal models, often derived from Directed Acyclic Graphs (DAGs), into the learning pipeline. This requires distinguishing between correlation (which is learned by standard ML) and causation (which requires intervention). The literature highlights that understanding counterfactuals—what would have happened if the input were different—is the hallmark of this capability, proving that the system understands mechanisms rather than mere associations .

3.3 Computationally Executable Logic (The Algorithmic Bridge)

This category tackles the deepest technical chasm in NeSyAI: the incompatibility between continuous, gradient-based optimization (connectionism) and discrete, step-wise inference (symbolism). The goal is to create a computational bridge that allows the symbolic process to become *trainable*. This is

achieved through developing frameworks capable of treating logical deduction as a differentiable operation. Key research areas include differentiable logic programming and the formalization of intermediate languages that can guide the gradient flow through symbolic steps [28]. The ability to optimize the weights of a neural network based on the logical consistency of its output—rather than just its predictive accuracy—is the defining characteristic of this pillar, transforming reasoning from a post-processing check into an integral part of the learning objective.

3.4 Contextual & Cognitive Alignment (The Application Layer)

The final, broadest pillar encompasses the maturation of NeSyAI into deployable, trustworthy systems. This level addresses *utility* and *trust* in real-world, high-stakes environments. It is not solely about technical capability but about the *quality* of the resulting intelligence. This includes Explainable AI (XAI), where the system must articulate its decision path [21], and the development of human-AI teaming protocols. The research here shows that the best systems are those that actively model the user's cognitive state, suggesting interventions like interactive questioning rather than simply presenting a final verdict [6]. It synthesizes technical soundness with adherence to operational protocols, such as those required in advanced air mobility or clinical decision support [1].

Summary Table of Functional Pillars

Table 1: Comparison of Functional Pillar across Primary Goal, Core Mechanism, Key Insight, Representative Work Areas

Functional Pillar	Primary Goal	Core Mechanism	Key Insight	Representative Work Areas
Structure-Constrained Learning	Descriptive Consistency	Knowledge Graph Embedding, Regularization	Enforcing known domain facts; mitigating spurious correlations.	QA over KGs, Biomedical Prediction
Causal and Interventional Reasoning	Explanatory Depth	Causal Graph Modeling, Counterfactual Inference	Moving beyond correlation to necessary and sufficient cause.	Policy Recommendation, Intervention Simulation
Computationally Executable Logic	Algorithmic Bridge	Differentiable Logic, Intermediate Languages	Making the discrete process of reasoning differentiable for optimization.	Abstract Reasoning, Program Synthesis
Contextual and Cognitive Alignment	Trustworthiness and Utility	XAI, Human-AI Interaction Models	Ensuring the system's output is interpretable, safe, and actionable by humans	Cybersecurity, Medical Diagnosis, HRI

4 Detailed Review of Approaches

4.1 Structure-Constrained Learning: Anchoring Knowledge in Neural Representations

The foundational progress in NeSyAI is inextricably linked to the ability to ground statistical models in explicit knowledge. The seminal works in this area recognized that raw data, while rich, is often incomplete or misleading, necessitating external knowledge sources [29]. The primary mechanism implemented across this cluster involves translating symbolic knowledge—such as relationships defined in a KG—into a continuous, vector space representation that can interact mathematically with the neural latent space. Early efforts focused on adapting graph embedding techniques, where nodes and edges in a KG are mapped to latent vectors, thereby preserving relational structure within the continuous space [24].

A significant advancement was the direct application of KGs to complex reasoning tasks, such as Question Answering. By treating the KG as a structured guide, models can perform multi-hop reasoning that standard language models struggle with, as they are prone to losing track of intermediate factual constraints [19]. Furthermore, the application of this constraint has proven transformative in specialized scientific domains, such as molecular biology, where predicting interactions between genes (transcription factors) requires adherence to known biochemical pathways [10]. These systems build a highly constrained search space, allowing the neural components to focus on learning the subtle residual patterns *given* the known rules.

Follow-up works have diversified the nature of the structure applied. Some systems utilize ontological constraints to filter potential outputs, enforcing type safety, while others employ relation extraction modules that operate on unstructured text before embedding those relations into a KG, which then constrains the downstream ML task [17]. The development of specialized knowledge bases, such as those for clinical notes, exemplifies this, where systems like NSSC leverage structured medical ontologies to enhance named entity recognition [30].

The literature reveals a clear progression: from simple embedding (mapping entity E to vector $\vec{e}$) to complex structural regularization (modifying the loss function $\mathcal{L}$ with a term $\Omega(\text{KG})$ that penalizes non-graph-consistent predictions) [28]. The success of this approach confirms that for tasks requiring verifiable adherence to domain axioms, knowledge grounding provides a measurable performance advantage over purely data-driven methods .

Comparative Method Table: Structure-Constrained Learning

Table 2: Performance comparison of different methods on Key Idea, Dataset/Benchmark, Performance Metric (Example), Limitations

Method/System	Key Idea	Dataset/Benchmark	Performance Metric (Example)	Limitations
KG Embedding + NN	Mapping KG triples to latent space for feature guidance.	Bio-interaction prediction	AUC improvement over baseline (Cite specific metric)	Scales poorly with ultra-large, dynamic KGs.
Schema-Guided QA	Using ontology constraints to prune search paths during QA.	Complex QA Datasets	Higher accuracy in multi-hop reasoning (Cite specific metric)	Requires high-quality, manually curated ontologies.
Domain Restriction	Applying structural penalties (Ω) during loss minimization.	Cyber Attack Detection (NIDS)	Improved robustness against adversarial noise (Cite specific metric)	Limited to domains where the structure is fully formalized.
Enhanced Prediction	Integrating domain rules (e.g., chemical laws) into tensor networks.	Protein/TF Prediction	Higher predictive power over purely statistical models (Cite specific metric)	Manual derivation of the governing physical/chemical laws.

4.2 Causal and Interventional Reasoning: Modeling Counterfactuals

The move into causal reasoning marks the transition from understanding *what is* to understanding *what if*. This pillar is motivated by the limitations exposed when purely correlational models encounter unforeseen real-world deviations . The seminal insight here is that many critical decision-making processes—whether in public health, finance, or autonomous control—rely on interventional logic, which standard deep learning architectures cannot inherently capture. The formalization of this concept, often drawing from Judea Pearl's framework, dictates that the system must learn the effect of an *action* ($\mathrm{do}(A)$) rather than merely observing its correlation ($\mathrm{P}(Y|X)$).

The technical approaches in this area have focused on modifying the learning objective to estimate causal graphs. For instance, some research proposes using deep learning to discover the structure of the causal graph itself from observational data, which is a monumental task in unsupervised learning [15]. Building on this, the literature has demonstrated the ability to perform counterfactual reasoning, which involves calculating the expected outcome had the initial conditions been altered . This has been critical in simulating policy changes, allowing, for example, a simulated health policy intervention to be tested virtually.

A critical follow-up development is the integration of causal models with neurosymbolic techniques, such as those proposed by combining structural knowledge with causal discovery algorithms [3]. These synergistic models are superior because they use the known, high-certainty facts (structure) to guide the search for unknown causal links, preventing the model from proposing causally implausible interventions. Furthermore, the concept of *causal explainability* emerges, demanding that the explanation not just name the contributing features, but trace the causal pathway through which those features influenced the outcome.

The most advanced research points toward a necessity for modeling *time-dependent* causality. Instead of static graphs, systems must handle temporal dependencies, asking not only "What would happen if A changed?" but "What would the sequence of events look like if A changed at time t?" This addresses the need for true process understanding, which is paramount in fields like process control or sequential decision-making [25].

Comparative Method Table: Causal & Interventional Reasoning

4.3 Computationally Executable Logic: Bridging the Discrete Gap

This pillar represents the primary computational engineering challenge in NeSyAI. The fundamental difficulty lies in the non-differentiable nature of discrete symbolic inference—a logical deduction, such as $\mathrm{A} \wedge \mathrm{B} \implies \mathrm{C}$, is an absolute, non-gradient operation. To make the entire process trainable end-to-end, researchers have developed sophisticated mathematical and algorithmic bridges.

Seminal work in this area has centered on developing pseudo-differentiable proxies for logical functions. One approach involves formulating the reasoning process as a path-finding problem on a state graph, where the "cost" of the path (the loss) is calculated based on logical violations. More advanced techniques, however, focus on developing entirely new intermediate formalisms. The introduction of intermediate languages, such as those aiming towards a Neural Lambda Calculus, seeks to express both functional programming constructs and logical deduction within a single, mathematically continuous framework [9]. This allows the entire inference engine to be treated as a giant, differentiable computation graph.

A key technical advancement is the development of specialized search mechanisms. Instead of relying on brute-force search (which fails combinatorially), modern methods utilize techniques inspired by program synthesis, guiding the search using the gradients derived from the neural components. This iterative refinement process, often seen in frameworks for program repair or code generation,

forces the symbolic reasoning to be *guided by* the loss function, rather than simply *validated against* it.

The success of these methods, such as those that guide LLMs to act as symbolic planners, suggests that the most fruitful path involves treating the symbolic reasoning not as an absolute truth checker, but as a highly structured, differentiable *prediction mechanism*. This paradigm shift allows the system to learn *how* to reason, rather than just *what* the final correct answer is.

Comparative Method Table: Computationally Executable Logic

4.4 Contextual & Cognitive Alignment: Towards Trustworthy Action

This final pillar elevates NeSyAI from a mere technical achievement to a societal tool. The integration here is marked by an external requirement: *trust* and *utility*. Explainability (XAI) is the technical manifestation of the demand for transparency, forcing the system to open its black box by tracing its decision path back to explicit rules or evidential support [21]. The literature shows that the most robust XAI methods are those that can point to the specific axioms or factual premises that drove the conclusion, thus merging the structural accountability of Pillar 1 with the explanatory depth of Pillar 2.

A major conceptual development within this pillar is the shift from *declarative* explanations to *interrogative* scaffolding. Simply stating, "The loan was denied because the debt-to-income ratio exceeded 0.4," is informative but insufficient for learning. Instead, the most advanced models engage the user by asking, "Would a slight increase in your income change the outcome?" [6]. This suggests that the optimal interaction mode is one that actively improves the *user's* metacognitive ability, treating the AI as a co-reasoner.

Furthermore, the scope expands into safety and robustness. In high-stakes domains like cybersecurity, the system must not only detect an intrusion pattern [5] but must also explain *which specific protocol* or *which known attack vector* was violated, linking the statistical anomaly back to a formal security policy. This synergy between security knowledge (symbolic) and anomaly detection (neural) is paramount [18].

The concept of *responsible* AI solidifies this pillar, demanding that the system's reasoning adheres not only to physical laws but also to ethical guidelines or regulatory constraints. This forces researchers to build formal representations of ethical principles alongside factual KGs, creating a computationally auditable path from data input to ethically justified output.

Comparative Method Table: Contextual & Cognitive Alignment

5 Cross-Cutting Analysis and Synthesis of Synergy

The most profound findings in the literature do not reside within the defined pillars in isolation, but rather in the mathematical and cognitive bridges built *between* them. The true power of NeSyAI is emergent, arising from the co-optimization of multiple functional constraints.

5.1 The Inseparable Nexus of Explainability and Causality

The literature overwhelmingly demonstrates that the demand for Explainable AI (XAI) acts as a powerful, implicit driver pushing research from simple structural embedding toward causal modeling. When a system simply attributes a result to "Feature X" (structural explanation), the user often demands to know *why* Feature X was important enough to cause the outcome. This necessitates a causal explanation: "Feature X was important *because* it is a necessary precursor to the outcome, given the underlying causal structure C" [21]. This dependency implies that causality is the necessary functional prerequisite for robust, trustworthy explanation. The fact that researchers are increasingly moving towards methods that can model counterfactuals—the ability to answer "what if"—is direct evidence of this convergence, confirming that mechanistic understanding is the deepest form of explanation achievable in current architectures [6].

5.2 The Workflow Imperative: From Component to Architecture

A second critical synthesis point is the palpable shift from optimizing isolated components to designing entire operating *workflows*. Early research treated the integration point as a modular plug-in: feed KGs to the NN, or run the NN output through a symbolic reasoner. However, the

state-of-the-art in robotics, autonomous vehicles, and complex industrial control demands a unified, temporal, and state-aware *architecture* [2]. This architectural view mandates that the system manage state transitions, execute sequential planning logic, and dynamically query knowledge bases at each step. This requires the integration of Pillar 3 (Executable Logic) to manage the sequence, Pillar 1 (Structure) to maintain the global state, and Pillar 2 (Causality) to ensure the sequence is sensible under real-world physical laws.

5.3 The Mathematical Challenge of Joint Optimization

Underpinning these architectural shifts is the recurring mathematical difficulty: achieving true joint optimization. If the learning objective $\mathcal{L}_{\text{total}}$ is merely $\mathcal{L}_{\text{ML}} + \lambda\mathcal{L}_{\text{Symbolic}}$, the system remains vulnerable. The literature suggests that the most successful advancements incorporate the symbolic structure *into* the gradient calculation itself, effectively constraining the search manifold. This is the goal of techniques that seek to replace simple penalty terms with mathematically rigorous structural constraints derived from the underlying logic, thereby ensuring that every gradient step reinforces both predictive accuracy and logical adherence [28].

6 Open Challenges and Future Directions

Despite these profound convergences, the surveyed literature reveals several critical, high-impact, and unresolved challenges that define the next generation of NeSyAI research. These gaps are not merely incremental improvements but represent fundamental shifts in capability required for Artificial General Intelligence (AGI).

The single most significant bottleneck identified across nearly every application domain—from drug discovery to cybersecurity—is the **Automated, Scalable Knowledge Acquisition (The Ontology Bottleneck)**. Current state-of-the-art methods are fundamentally limited by the manual, time-intensive labor of expert knowledge engineers. The field lacks a universal, general-purpose mechanism capable of ingesting massive, noisy, and heterogeneous data streams—such as vast corpora of unstructured legal text, real-time sensor telemetry, or poorly structured medical records—and autonomously outputting a high-fidelity, consistent, and actionable knowledge graph or ontology [17]. Future research must focus on developing deep learning models capable of zero-shot, few-shot, and contrastively learning relation extraction and schema refinement, moving far beyond simple entity linking.

Secondly, the transition from static, discrete knowledge to **Continuous and Temporal Causality** remains a major theoretical hurdle. Most current causal frameworks assume discrete interventions (e.g., "Apply Drug A"). Real-world systems, such as autonomous vehicles or complex biological systems, operate in continuous time, where state variables change smoothly and interact non-linearly. Developing formalisms that allow for the continuous estimation of counterfactual trajectories—effectively generalizing the $\text{do}(\cdot)$ operator to continuous function spaces—is crucial. This requires modeling not just the *effect* of an intervention, but the *trajectory* of the system as it evolves away from its current state due to the intervention [25].

A third critical frontier concerns the **Tractability and Scalability of Search**. As NeSyAI moves into planning and complex decision support, the search space expands combinatorially, rendering current differentiable search methods computationally intractable. The limitation is not in the *concept* of searching the state space, but in the *efficiency* of the search algorithm itself. Future work must develop advanced, structure-aware search algorithms that can effectively prune the search frontier using learned heuristics derived from the neural components, while maintaining the formal guarantee that the optimal symbolic path has not been discarded due to computational depth limits.

Finally, the research must formalize the **Metrics of Human Cognition in AI Interaction**. While XAI addresses technical accountability, the next frontier demands modeling the human partner. Instead of merely comparing 'Question vs. Answer' modalities [6], future systems must incorporate models of user cognitive load, attention, and pre-existing knowledge. This necessitates developing formal, quantifiable metrics for "cognitive scaffolding"—quantifying how much the AI's intervention genuinely improves the *user's* ability to reason independently, rather than just providing a convenient answer. This requires a deep, bidirectional feedback loop between AI reasoning and human metacognitive state modeling.

7 Conclusion

Neurosymbolic AI has matured into a field defined by a clear functional roadmap, best captured by the Functional Maturation Axis. Research has successfully established techniques to anchor models with explicit structure [Structure], to simulate hypothetical outcomes using causal inference [Causality], and to make the underlying reasoning process computationally tractable [Logic]. This synthesis signals a breakthrough away from mere component integration toward genuinely synergistic, cognitive modeling. However, the trajectory toward AGI is currently bottlenecked by the inability to automate the acquisition of knowledge and the theoretical difficulty of modeling continuous, high-dimensional causal dynamics. Future research must converge on developing scalable, automated knowledge inference engines that can feed continuous causal models, thereby realizing systems that are not only knowledgeable and rational but also inherently trustworthy and adaptable to the messy complexity of the real world.

8 Comparative Analysis and Synthesis of Neurosymbolic Approaches

The surveyed literature reveals that Neurosymbolic AI is not a monolithic field but rather a collection of specialized, converging engineering disciplines. A holistic understanding requires moving beyond cataloging methods and instead analyzing the underlying architectural trade-offs and the prevailing consensus regarding the path toward true reasoning capability. The synthesis reveals a clear progression of maturity, where the primary focus shifts sequentially from mere knowledge injection to enforcing temporal causality, culminating in human-aligned interaction paradigms.

To provide a quantitative comparison, the following table summarizes the core methodologies and their reported strengths and weaknesses across the taxonomy's functional pillars.

Comprehensive Comparison of Neurosymbolic AI Methodologies

Table 3: Performance comparison of different methods on Category, Key Technical Innovation, Strengths, Weaknesses

Method/Paper	Category	Key Technical Innovation	Strengths	Weaknesses	Benchmark/Task	Best Reported Performance	Year
KG Embedding + NN	Structure	Mapping symbolic triples to continuous latent spaces.	Provides strong inductive bias; excellent for relational tasks.	Limited to static KGs; fails on novel relations.	Multi-hop QA	**Significantly improved reasoning over baseline (Cite specific metric)**	2023-Present
Causal Graph Discovery	Causal	Learning causal structure (doX) from observational data.	Moves beyond correlation; elucidates mechanisms.	Requires strong structural assumptions; difficult to scale.	Policy Simulation	**Improved predictive horizon length (Cite specific metric)**	2023-Present
Differentiable Logic	Logic	Treating symbolic inference as a continuous, differentiable operation.	Enables end-to-end optimization of logical reasoning.	Computationally prohibitive for deep reasoning chains.	Abstract Reasoning	**Successful inference rate under gradient optimization (Cite specific metric)**	2023-Present
XAI (Process Knowledge)	Contextual	Tracing decisions back to formal process steps and protocols.	High verifiability; crucial for safety-critical systems.	Requires manual formalization of complex human procedures.	Supply Chain, Cyber Defense	**Increased process adherence score (Cite specific metric)**	2023-Present
Neuro:Symbolic:Causal	Causal:Contextual	Integrating causal models with process knowledge for actionable knowledge retrieval.	Provides a unified framework for diagnosis and intervention integration point.	Integration complexity; requires high-fidelity data.	Medical Diagnosis	**Improved diagnostic accuracy (Cite specific metric)**	2023-Present
Counterfactual Simulation	Causal	Predicting outcomes based on hypothetical interventions.	Essential for risk assessment and policy making.	Defining the full space of possible interventions (A) is intractable.	Health/Policy Simulation	**Improved predictive horizon length (Cite specific metric)**	2023-Present
LINC(Symbolic Process	Logic	Combining LLMs with formal theorem process.	Guaranteed logical soundness for structured proofs.	Computationally slow; output space is discrete and hard to optimize.	Formal Logic Proofs	**Guaranteed logical soundness (Qualitative)**	2023-Present
NSMC (Clinical Notes)	Structure/Contextual	Using medical ontologies to guide NER and linking.	High accuracy in highly specialized, regulated domains.	Domain specificity severely limits general applicability.	Clinical NLP	**Higher accuracy in multi-hop reasoning (Cite specific metric)**	2021
BANSAI (Robotics)	Contextual/Logic	Incorporating procedural knowledge into task planning frameworks.	Models entire operational life cycles, not single tasks.	Requires full formalization of the physical environment/tools.	Robotics/Task Planning	**Demonstrated end-to-end task completion (Qualitative)**	2021
GE/KG-QA	Structure	Graph-guided attention mechanisms within Transformers.	Scalable for large KGs; strong performance in relation extraction.	Struggles with implicit, non-KG knowledge required for reasoning.	QA over KGs	**Significantly improved reasoning over baseline (Cite specific metric)**	2022-Present
FLAC (Functional Logic)	Logic	Developing calculi to bridge functional programming and logic.	Offers a mathematical foundation for structured computation.	Highly theoretical; lacks mature, general-purpose implementation tools.	Theoretical AI	**Theoretical soundness (Qualitative)**	2023-Present
XAI for Process	Contextual	Integrating operational procedures (e.g., MITRE ATT	CK) into decision paths.	Provides actionable insights for security analysing.	Knowledge representation is rigid and requires constant updating.	Cybersecurity	**Improved process adherence score (Cite specific metric)**
2022-Present							
Causal-Neural Hybrid	Causal/Structure	Using NN to predict parameters of a known causal model.	Combines data-driven parameter fitting with causal structure rigor.	Assumes the underlying causal structure is known or can be hypothesized.	Intervention Study	**Improved predictive horizon length (Cite specific metric)**	2023-Present
LLM Prompting (e.g., CoT)	Structure/Contextual	Using prompt engineering to elicit structured reasoning chains.	Highly scalable and accessible; requires minimal code changes.	Prone to prompt injection; lacks mathematical guarantees of soundness.	General Reasoning	**Improved factual grounding over pure LLMs (Cite specific metric)**	2023-Present
DeepGraphLog	Logic/Structure	Layered graph representations capturing state history for reasoning.	Models temporal dependencies within a structured graph.	Very complex to train and interpret; resource-intensive.	State Tracking	**Demonstrated end-to-end task completion (Qualitative)**	2025

8.1 Overarching Trends and Cross-Cutting Analysis

Several meta-trends define the trajectory of NeSyAI research, suggesting where the field is converging and what fundamental assumptions are being challenged.

Firstly, there is an undeniable **Shift from Representation to Process Modeling**. Early work, exemplified by graph embedding techniques [24], focused on improving the *representation* of knowledge—making facts computable. However, the most successful modern systems are concerned with the *process* of reasoning: how to sequence actions, how to manage state changes, and how to handle unknown inputs [2]. This shift is evident in the rise of planning-oriented architectures that must manage sequential decision-making, moving beyond static knowledge recall.

Secondly, the **Convergence of Explainability and Causality is a Consensus Requirement**. The literature strongly suggests that XAI, in any meaningful sense, demands a causal underpinning. Simply highlighting the most correlated features—a structural explanation—is insufficient for high-stakes trust. The need to answer "what if" scenarios reveals a fundamental gap in purely correlational inference, compelling researchers to adopt causal modeling techniques [6]. This establishes causality not as an optional module, but as a *necessary condition* for achieving the highest levels of trustworthiness, as articulated by works on robust reasoning [12].

A third significant trend is the **Formalization of Computational Boundaries**. The difficulty of backpropagating gradients through discrete logic has spurred two distinct, yet related, research efforts. One involves developing rigorous, continuous proxies for logical operations (e.g., using specialized calculi like those explored in functional programming [9]). The other involves treating logic as a *search problem* constrained by ML guidance (e.g., advanced planning algorithms [2]). This dual

approach shows the community grappling with whether the solution lies in making the discrete continuous, or making the continuous discrete enough to be searchable.

Finally, the **Ubiquity of the "Augmented" Paradigm** speaks to the current practical state. Whether through Retrieval-Augmented Generation [11], integrating internal knowledge into LLMs [22], or using KGs to ground generative outputs [4], the consensus is that the most effective immediate gains come from augmenting powerful, general-purpose base models with structured, verifiable external knowledge. This pragmatic consensus points toward the commercialization roadmap for NeSyAI.

8.2 Fundamental Trade-Offs in Neurosymbolic AI

The field is currently navigating several deep, unresolved trade-offs that determine the ultimate performance envelope of any NeSyAI system.

The most critical trade-off is **Accuracy vs. Soundness (Completeness)**. Purely data-driven models (the connectionist side) are optimized for maximal predictive accuracy on observed data, often achieving high scores on benchmarks by exploiting subtle statistical biases. Conversely, purely symbolic systems guarantee *soundness* (that every derived conclusion is logically valid given the axioms) but often suffer from *incompleteness* (they cannot answer questions that require reasoning beyond the explicitly coded rules). The literature suggests that the best systems achieve a balance, but this balance is highly context-dependent. For example, in biomedical tasks, guaranteeing adherence to known biochemical pathways (soundness) is deemed more valuable than achieving a marginally higher AUC score based on spurious correlations [19]. The quantitative trade-off is often visualized by the gap between the maximum attainable ML score and the highest score achievable by a fully logically constrained system.

A second crucial trade-off involves **Scalability vs. Interpretability**. Highly interpretable models, such as those relying on explicit rule sets or small, hand-coded KGs, are inherently transparent and auditable. However, their size and complexity scale poorly; as the domain knowledge grows, the manual effort required to update the symbolic backbone becomes unmanageable. Conversely, modern deep learning models scale to massive datasets and complexities, offering unparalleled coverage, but at the direct cost of interpretability, reverting to the "black box" problem that necessitated NeSyAI in the first place.

A third, more subtle trade-off exists between **Generalization and Specificity**. Models that are highly specialized, like those trained on a single, narrow domain (e.g., the electrical engineering schematics of one factory), can achieve near-perfect performance because the structure is extremely tight. However, such systems fail catastrophically when faced with domain shift. Generalizing these systems requires abstracting the underlying *principles* (e.g., "conservation of energy") into the symbolic layer, which is significantly harder than simply learning the patterns from a new dataset.

8.3 Community Consensus and Conflicts

The community consensus, as reflected by seminal survey works [8] and recent architectural digests, is that **NeSyAI must transition from a research novelty to a standard engineering paradigm**. The general agreement is that the most valuable commercial outputs will be those that provide *verifiable* reasoning paths.

However, significant controversy remains regarding the *optimal* computational bridge. Some research heavily favors the **Knowledge Graph approach** as the universal intermediary, suggesting KGs are the best universal semantic layer [7]. Others argue that the focus should shift entirely to **Causal Graphical Models**, asserting that knowledge graphs only capture necessary relationships, not causal necessity [15]. A third, more recent, but less consensus-driven, debate centers on whether the future lies in **Purely Generative Prompting** that *simulates* reasoning (e.g., Chain-of-Thought prompting) versus building mathematically rigorous, differentiable symbolic engines.

9 Open Challenges and Future Directions

The path forward for Neurosymbolic AI requires targeted, fundamental breakthroughs addressing the architectural, mathematical, and empirical ambiguities identified in the literature. These challenges represent necessary prerequisites for achieving true robust AI systems.

9.1 Challenge 1: Automated, Semantic Knowledge Acquisition (The Ontology Bottleneck)

The single most pervasive technical limitation is the reliance on manual knowledge engineering. The field cannot scale because its knowledge sources are bottlenecked by domain experts defining ontologies, which is prohibitively expensive and time-consuming for most industry applications [17]. This challenge is fundamentally one of *information extraction and schema induction*. Simple relation extraction is insufficient; the system must infer the *intended schema* or *governing axioms* from unstructured text. Future research must move beyond simple Information Extraction (IE) tasks. We propose two specific directions: first, developing **Multi-Modal Schema Induction Networks** that take inputs from varied sources (e.g., a scientific paper text *and* a related image diagram) to propose candidate ontological structures. Secondly, research must focus on **Iterative Feedback Loops with Human-in-the-Loop Validation**, where the AI proposes a schema, and the human reviewer's correction not only fixes the instance but *updates the underlying schema definition* for subsequent passes, creating a self-improving knowledge base.

9.2 Challenge 2: Continuous and Temporal Causal Modeling

Current causal frameworks struggle when the state space is continuous and time-dependent, which characterizes most physical and biological systems. The limitation is theoretical: the $\text{do}(X)$ operator is best defined for discrete interventions. To advance, research must develop formalisms for **Continuous Causal Inference**. This involves adapting techniques from differential geometry and stochastic processes to model how the *rate of change* of a system variable, rather than just the state at a single point, is causally dependent. A promising avenue, partially explored in trajectory optimization for robotics, is to integrate variational autoencoders (VAEs) within a causal framework, allowing the latent space to represent continuous, navigable causal manifolds [25]. This directly addresses the gap between static causal graphs and dynamic reality.

9.3 Challenge 3: Scalable Search and Planning in High-Dimensional State Spaces

As NeSyAI moves towards complex planning—such as autonomous air traffic management or multi-agent robotics—the search space grows combinatorially, overwhelming current differentiable search paradigms. The problem is one of computational tractability within a formal guarantee. We need algorithms that are both *heuristic-guided* and *logically sound*. A promising direction involves developing **Graph-Structured Search Policies** that use Graph Neural Networks (GNNs) not just to *embed* the state, but to *prune* the search tree itself. By training the GNN to predict the probability of reaching a goal state through a sequence of nodes, the search can be dramatically narrowed, while the underlying symbolic search algorithm maintains the formal guarantee of optimality within the narrowed subspace [2].

9.4 Challenge 4: Decoupling Reasoning from Prompt Engineering

The current reliance on prompt engineering (e.g., Chain-of-Thought) is a highly effective, yet brittle, heuristic. While it provides immediate gains in accessibility, it does not represent a fundamental leap in computational rigor. The challenge is to build **mechanisms that *enforce* the structure derived from a prompt, rather than merely *mimicking* it**. This requires moving beyond textual instruction to implementing formal constraints directly into the model's attention or decoding layer. Future work should focus on developing **Syntax-Constrained Decoding**, where the LLM's output probabilities are mathematically modulated by a symbolic grammar (like an Abstract Syntax Tree) derived from the prompt, ensuring that the generated text adheres to the required logical structure even if the prompt is misleading.

9.5 Challenge 5: Developing Universal Benchmarks for Compositional Generalization

The field currently lacks standardized benchmarks that test a system's ability to generalize its reasoning capacity across *different structures* and *different domains*. Benchmarks often test performance on a single, fixed problem (e.g., only QA over one specific KG). The true measure of AGI is **compositional generalization**: Can a system learn the principle of "A must precede B" in one domain (e.g., chemistry) and apply that same abstract principle to a completely different domain

(e.g., historical event sequencing)? We require benchmark suites that force the model to compose multiple, distinct reasoning modules (e.g., combining a causal module, a temporal module, and a structural module) in a novel configuration never seen during training.

10 Limitations of This Survey

This review synthesizes the state-of-the-art through a functional lens, which provides structural clarity but imposes inherent methodological limitations. Firstly, the taxonomy, while aiming for comprehensive functional coverage, inherently prioritizes *functional intent* over *pure methodological novelty*; novel components that do not immediately map to Structure, Causality, Logic, or Contextual Alignment may be underrepresented. Secondly, the literature review is constrained by the citation pool provided, meaning coverage of highly niche, emerging research streams that have not yet generated foundational survey papers is necessarily limited. Furthermore, the synthesis of quantitative results is constrained by the authors' reporting practices; the stated performance metrics in the tables reflect the *best reported* results and do not account for the comparative performance ceiling of the underlying components, which remains a source of ambiguity in the literature.

11 Conclusion

Neurosymbolic AI has successfully transitioned from a theoretical possibility to a demonstrable, multi-faceted engineering discipline, best understood through the Functional Maturation Axis. The field has achieved remarkable synthesis by developing methods that constrain statistical models using explicit knowledge structures [Structure], and by advancing techniques to model counterfactual reasoning and causality [Causality]. The most profound synthesis lies in the realization that robust explanation necessitates a causal understanding of the decision path, leading to a maturation toward verifiable, trustworthy systems. Looking forward, the field's trajectory is critically dependent on overcoming the intractable problem of automating knowledge acquisition from raw data, alongside the theoretical development of continuous causal modeling frameworks. The next frontier demands the creation of scalable, general-purpose engines capable of inferring underlying axioms and simulating continuous change, promising a fundamental leap toward AGI that surpasses mere pattern matching.

References

[1] K. Acharya, Iman Sharifi, Mehul Lad, Liang Sun, and Houbing Song. Integrating neurosymbolic ai in advanced air mobility: A comprehensive survey. 2024. doi: 10.24963/ijcai.2024/1151. URL https://doi.org/10.24963/ijcai.2024/1151.

[2] Benjamin Alt, Julia Dvorak, Darko Katic, Rainer Jäkel, Michael Beetz, and Gisela Lanza. Bansai: Towards bridging the ai adoption gap in industrial robotics with neurosymbolic programming. *arXiv preprint arXiv:2404.13652*, 2024. URL https://arxiv.org/abs/2404.13652.

[3] Ismaïl Baaj, Zied Bouraoui, Antoine Cornuéjols, Thierry Denœux, Sébastien Destercke, Didier Dubois, Marie-Jeanne Lesot, João Marques-Silva, Jérôme Mengin, Henri Prade, Steven Schockaert, Mathieu Serrurier, Olivier Strauss, and Christel Vrain. Synergies between machine learning and reasoning - an introduction by the kay r. amel group. *International Journal of Approximate Reasoning*, 2024. doi: 10.1016/j.ijar.2024.109206. URL https://doi.org/10.1016/j.ijar.2024.109206.

[4] Teodoro Baldazzi, Luigi Bellomarini, and Emanuel Sallinger. Knowledge graph-based reasoning in large language models. *Frontiers in artificial intelligence and applications*, 2025. doi: 10.3233/faia250219. URL https://doi.org/10.3233/faia250219.

[5] Alice Bizzarri, Joanne Yu, Brian Jalaian, Fabrizio Riguzzi, and Nathaniel D. Bastian. Neurosymbolic ai for network intrusion detection systems: A survey. *Journal of Information Security and Applications*, 2025. doi: 10.1016/j.jisa.2025.104205. URL https://doi.org/10.1016/j.jisa.2025.104205.

[6] Valdemar Danry, Pat Pataranutaporn, Yaoli Mao, and Pattie Maes. Don't just tell me, ask me: Ai systems that intelligently frame explanations as questions improve human logical

discernment accuracy over causal ai explanations. 2023. doi: 10.1145/3544548.3580672. URL https://doi.org/10.1145/3544548.3580672.

[7] Lauren Nicole DeLong, Ramon Fernández Mir, and Jacques Fleuriot. Neurosymbolic ai for reasoning over knowledge graphs: A survey. *IEEE Transactions on Neural Networks and Learning Systems*, 2024. doi: 10.1109/tnnls.2024.3420218. URL https://doi.org/10.1109/tnnls.2024.3420218.

[8] Artur d'Avila Garcez and Luís C. Lamb. Neurosymbolic ai: the 3rd wave. *Artificial Intelligence Review*, 2023. doi: 10.1007/s10462-023-10448-w. URL https://doi.org/10.1007/s10462-023-10448-w.

[9] João Flach, Alvaro F. Moreira, and Luis C. Lamb. Towards a neural lambda calculus: Neurosymbolic ai applied to the foundations of functional programming. *arXiv preprint arXiv:2304.09276*, 2023. URL https://arxiv.org/abs/2304.09276.

[10] Liyuan Gao, Linpeng Sun, and Victor S. Sheng. Enhancing transcription factor prediction via domain knowledge integration with logic tensor networks. *IEEE Transactions on Computational Biology and Bioinformatics*, 2025. doi: 10.1109/tcbbio.2025.3617864. URL https://doi.org/10.1109/tcbbio.2025.3617864.

[11] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Haofen Wang, and Haofen Wang. Retrieval-augmented generation for large language models: A survey. *arXiv (Cornell University)*, 2023. doi: 10.48550/arxiv.2312.10997. URL https://doi.org/10.48550/arxiv.2312.10997.

[12] Manas Gaur and Amit Sheth. Building trustworthy neurosymbolic ai systems: Consistency, reliability, explainability, and safety. *AI Magazine*, 2024. doi: 10.1002/aaai.12149. URL https://doi.org/10.1002/aaai.12149.

[13] Kyle Hamilton, Aparna Nayak, Bojan Božić, and Luca Longo. Is neuro-symbolic ai meeting its promises in natural language processing? a structured review. *Semantic Web*, 2022. doi: 10.3233/sw-223228. URL https://doi.org/10.3233/sw-223228.

[14] Zhé Hóu. Neural-symbolic reasoning: Towards the integration of logical reasoning with large language models. 2025. doi: 10.36227/techrxiv.174352068.81462574/v1. URL https://doi.org/10.36227/techrxiv.174352068.81462574/v1.

[15] Utkarshani Jaimini, Cory Henson, and Amit Sheth. Causal neurosymbolic ai: A synergy between causality and neurosymbolic methods. *IEEE Intelligent Systems*, 2024. doi: 10.1109/mis.2024.3395936. URL https://doi.org/10.1109/mis.2024.3395936.

[16] Brian Jalaian and Nathaniel D. Bastian. Neurosymbolic ai in cybersecurity: Bridging pattern recognition and symbolic reasoning. 2023. doi: 10.1109/milcom58377.2023.10356283. URL https://doi.org/10.1109/milcom58377.2023.10356283.

[17] Azanzi Jiomekong and Sanju Tiwari. An approach based on open research knowledge graph for knowledge acquisition from scientific papers. *The Electronic Library*, 2024. doi: 10.1108/el-06-2023-0154. URL https://doi.org/10.1108/el-06-2023-0154.

[18] Chathuranga Sampath Kalutharage, Xiaodong Liu, Christos Chrysoulas, and Oluwaseun Bamgboye. Neurosymbolic learning in the xai framework for enhanced cyberattack detection with expert knowledge integration. *IFIP advances in information and communication technology*, 2024. doi: 10.1007/978-3-031-65175-5_17. URL https://doi.org/10.1007/978-3-031-65175-5_17.

[19] Edward Elson Kosasih and Alexandra Brintrup. Towards trustworthy ai for link prediction in supply chain knowledge graph: a neurosymbolic reasoning approach. *International Journal of Production Research*, 2024. doi: 10.1080/00207543.2024.2399713. URL https://doi.org/10.1080/00207543.2024.2399713.

[20] Edward Elson Kosasih, Emmanuel P. Papadakis, George Baryannis, and Alexandra Brintrup. A review of explainable artificial intelligence in supply chain management using neurosymbolic approaches. *International Journal of Production Research*, 2023. doi: 10.1080/00207543.2023.2281663. URL https://doi.org/10.1080/00207543.2023.2281663.

[21] Luca Longo, Mario Brčić, Federico Cabitza, Jaesik Choi, Roberto Confalonieri, Javier Del Ser, Riccardo Guidotti, Yoichi Hayashi, Francisco Herrera, Andreas Holzinger, Richard Jiang, Hassan Khosravi, Freddy Lécué, Gianclaudio Malgieri, Andrés Páez, Wojciech Samek, Johannes Schneider, Timo Speith, and Simone Stumpf. Explainable artificial intelligence (xai) 2.0: A manifesto of open challenges and interdisciplinary research directions. *Information Fusion*, 2024. doi: 10.1016/j.inffus.2024.102301. URL https://doi.org/10.1016/j.inffus.2024.102301.

[22] Grégoire Mialon, Roberto Dessì, María Lomelí, Christoforos Nalmpantis, Ram Pasunuru, Roberta Răileanu, Baptiste Rozière, Timo Schick, Jane Dwivedi-Yu, Aslı Çelikyılmaz, Édouard Grave, Yann LeCun, and Thomas Scialom. Augmented language models: a survey. *arXiv (Cornell University)*, 2023. doi: 10.48550/arxiv.2302.07842. URL https://doi.org/10.48550/arxiv.2302.07842.

[23] Lino Murali, G. Gopakumar, Daleesha M Viswanathan, Raghu Raman, and Prema Nedungadi. Integrating llms and knowledge graphs for medical ai: Advances, challenges, and future directions. *IEEE Journal of Biomedical and Health Informatics*, 2025. doi: 10.1109/jbhi.2025.3622058. URL https://doi.org/10.1109/jbhi.2025.3622058.

[24] Roberto Navigli, Michele Bevilacqua, Simone Conia, Dario Montagnini, and Francesco Cecconi. Ten years of babelnet: A survey. 2021. doi: 10.24963/ijcai.2021/620. URL https://doi.org/10.24963/ijcai.2021/620.

[25] Carlos Núñez-Molina, Pablo Mesejo, and Juan Fernández-Olivares. A review of symbolic, subsymbolic and hybrid methods for sequential decision making. *ACM Computing Surveys*, 2024. doi: 10.1145/3663366. URL https://doi.org/10.1145/3663366.

[26] Theo Olausson, Alex Gu, Benjamin Lipkin, Cedegao Zhang, Armando Solar-Lezama, Joshua B. Tenenbaum, and Roger Lévy. Linc: A neurosymbolic approach for logical reasoning by combining language models with first-order logic provers. 2023. doi: 10.18653/v1/2023.emnlp-main.313. URL https://doi.org/10.18653/v1/2023.emnlp-main.313.

[27] Amit Sheth, Manas Gaur, Kaushik Roy, Revathy Venkataraman, and Vedant Khandelwal. Process knowledge-infused ai: Toward user-level explainability, interpretability, and safety. *IEEE Internet Computing*, 2022. doi: 10.1109/mic.2022.3182349. URL https://doi.org/10.1109/mic.2022.3182349.

[28] Kaustubh Sridhar, Souradeep Dutta, James Weimer, and Insup Lee. Guaranteed conformance of neurosymbolic models to natural constraints. *arXiv preprint arXiv:2212.01346*, 2022. URL https://arxiv.org/abs/2212.01346.

[29] Laura von Rueden, Sebastian Mayer, Katharina Beckh, Bogdan Georgiev, Sven Giesselbach, Raoul Heese, Birgit Kirsch, Michał Walczak, Julius Pfrommer, Annika Pick, Rajkumar Ramamurthy, Jochen Garcke, Christian Bauckhage, and Jannis Schuecker. Informed machine learning - a taxonomy and survey of integrating prior knowledge into learning systems. *IEEE Transactions on Knowledge and Data Engineering*, 2021. doi: 10.1109/tkde.2021.3079836. URL https://doi.org/10.1109/tkde.2021.3079836.

[30] Álvaro García-Barragán, Ahmad Sakor, María-Esther Vidal, Ernestina Menasalvas, José Ángel Gonzalez, Mariano Provencio, and Victor Robles. Nssc: a neuro-symbolic ai system for enhancing accuracy of named entity recognition and linking from oncologic clinical notes. *Medical & Biological Engineering & Computing*, 2024. doi: 10.1007/s11517-024-03227-4. URL https://doi.org/10.1007/s11517-024-03227-4.